Resources and Transactions

a fundamental duality in observability

QCon Plus

Track: “Observability and Understandability in Production”

May 18, 2021

Ben Sigelman

Co-founder and CEO, Lightstep

Co-creator of OpenTelemetry and OpenTracing

Co-creator of Dapper and Monarch at Google

Part I

Transactions

What are “Transactions?”

- The things in your app that **actually “do something”** for the end-user
- These days, they’re distributed
- Can be described at *many* granularities… For example:
  - End-user workflows in mobile apps, etc
  - HTTP requests in microservices
  - Local function calls
  - CPU instructions

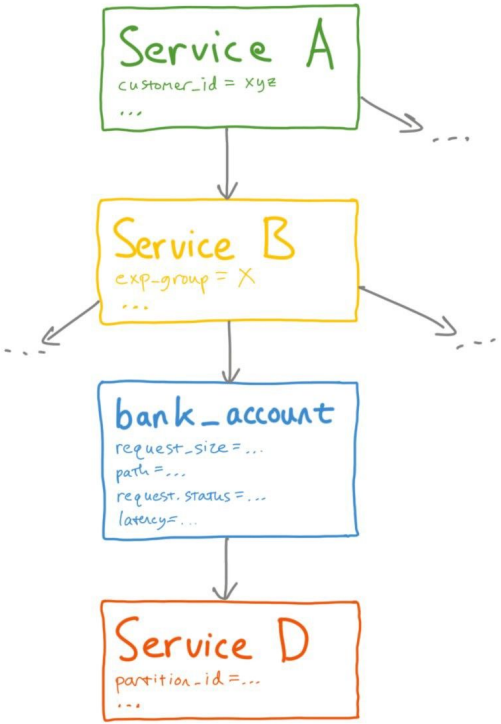

What are “Transactions?”

Telemetry: tracing and structured logs, with key:value attributes.

... and **tracing will eventually replace logging.**

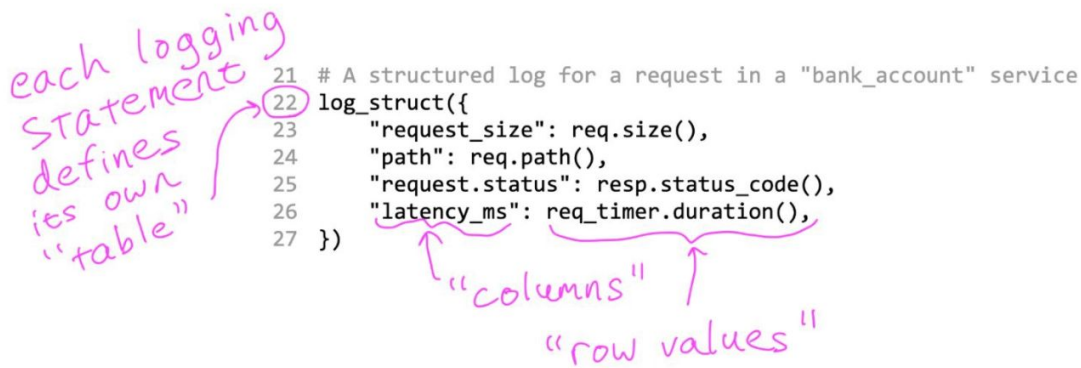

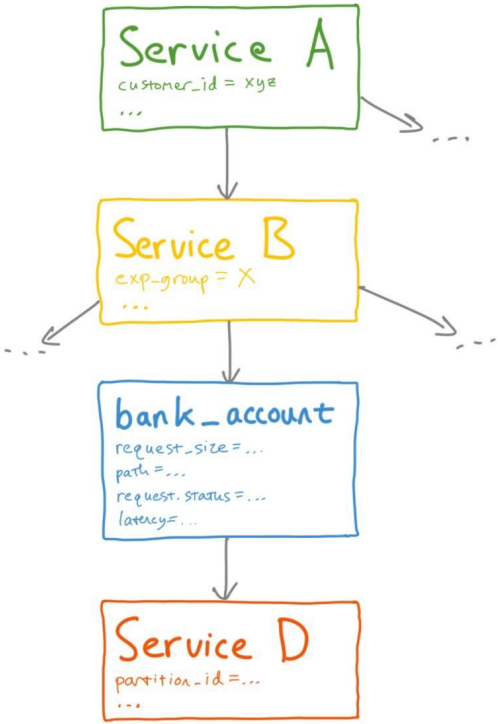

Tracing is just a JOIN across transactions!

relational table for the `bank_account (line 22)` log statement

explicit fields			
request_size	path	request.status	latency_ms
2641	/deposit	200	271
2573	/withdraw	200	198
2979	/embiggen	404	120
1265	/deposit	200	83
2325	/deposit	200	135
1799	/withdraw	200	392
1312	/deposit	200	122
1053	/withdraw	200	252
2844	/withdraw	503	397
2629	/withdraw	200	306
2549	/deposit	200	322
2270	/deposit	200	289

values from the log statement

fields from the log statement

each logging statement defines its own "table"

```
21 # A structured log for a request in a "bank_account" service
22 log_struct({
23     "request_size": req.size(),
24     "path": req.path(),
25     "request.status": resp.status_code(),
26     "latency_ms": req_timer.duration(),
27 })
```

"columns"

"row values"

Tracing is just a JOIN across transactions!

Implementing "JOIN" via distributed tracing

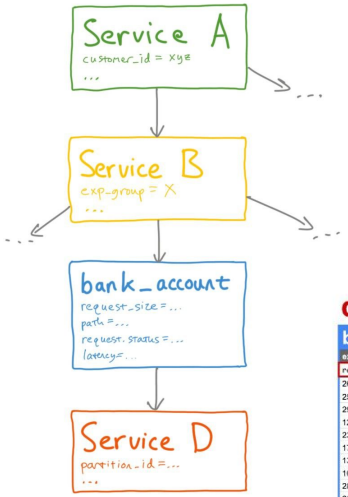

relational table for the `bank_account (line 22)` log statement

values from the log statement

fields from the log statement

explicit fields			
request_size	path	request.status	latency_ms
2641	/deposit	200	271
2573	/withdraw	200	198
2979	/embiggen	404	120
1265	/deposit	200	83
2325	/deposit	200	135
1799	/withdraw	200	392
1312	/deposit	200	122
1053	/withdraw	200	252
2844	/withdraw	503	397
2629	/withdraw	200	306
2549	/deposit	200	322
2270	/deposit	200	289

distributed tracing JOIN table for the `bank_account (line 22)` log statement

bank_account service							Service A				Service B				Service D			
explicit fields				implicit fields			explicit fields		implicit fields		explicit fields		implicit fields		explicit fields		implicit fields	
request_size	path	request.status	latency	version	hostname	timestamp	customer_id	...	version	...	exp_group	...	version	...	partition_id	...	version	...
2641	/deposit	200	271	3.5.24	ip42.ec2.internal	1618164815.298	2377	...	7.1.2	...	C	...	6.2.9	...	59	...	14.1.2	...
2573	/withdraw	200	198	3.5.24	ip42.ec2.internal	1618164845.965	5524	...	7.1.2	...	A	...	6.3.0	...	72	...	14.1.2	...
2979	/embiggen	404	120	3.5.24	ip79.ec2.internal	1618164865.558	6408	...	7.1.2	...	B	...	6.2.9	...	68	...	14.1.2	...
1265	/deposit	200	83	3.5.24	ip42.ec2.internal	1618164871.507	5900	...	7.1.2	...	A	...	6.2.9	...	67	...	14.1.2	...
2325	/deposit	200	135	3.5.24	ip42.ec2.internal	1618164878.719	3268	...	7.1.2	...	C	...	6.3.0	...	78	...	14.1.2	...
1799	/withdraw	200	392	3.5.24	ip79.ec2.internal	1618164959.044	3262	...	7.1.2	...	B	...	6.3.0	...	62	...	14.1.2	...
1312	/deposit	200	122	3.5.24	ip42.ec2.internal	1618164967.531	7746	...	7.1.2	...	B	...	6.3.0	...	43	...	14.1.2	...
1053	/withdraw	200	252	3.5.24	ip42.ec2.internal	1618165006.605	441	...	7.1.2	...	A	...	6.3.0	...	51	...	14.1.2	...
2844	/withdraw	503	397	3.5.25	ip54.ec2.internal	1618165046.479	3879	...	7.1.2	...	B	...	6.2.9	...	58	...	14.1.2	...
2629	/withdraw	200	306	3.5.24	ip79.ec2.internal	1618165048.559	4788	...	7.1.2	...	B	...	6.2.9	...	76	...	14.1.2	...
2549	/deposit	200	322	3.5.24	ip79.ec2.internal	1618165063.302	7697	...	7.1.2	...	A	...	6.2.9	...	59	...	14.1.2	...
2270	/deposit	200	289	3.5.24	ip42.ec2.internal	1618165067.863	1494	...	7.1.2	...	C	...	6.2.9	...	38	...	14.1.2	...

Tracing automates a system-wide JOIN!

Part II

Resources

What are “Resources?”

- **The things transactions *consume*** in order to do their work
  - Corollary: they are finite
- Resources **survive across transactions**
- Resources are **shared across transactions** (unless you have unlimited budget)
- Can also be described at *many* granularities... For example:
  - <Your AWS bill>
  - Kafka topic throughput
  - VM cpu usage or VM memory usage
  - Contention on a single mutex lock

What are resources?

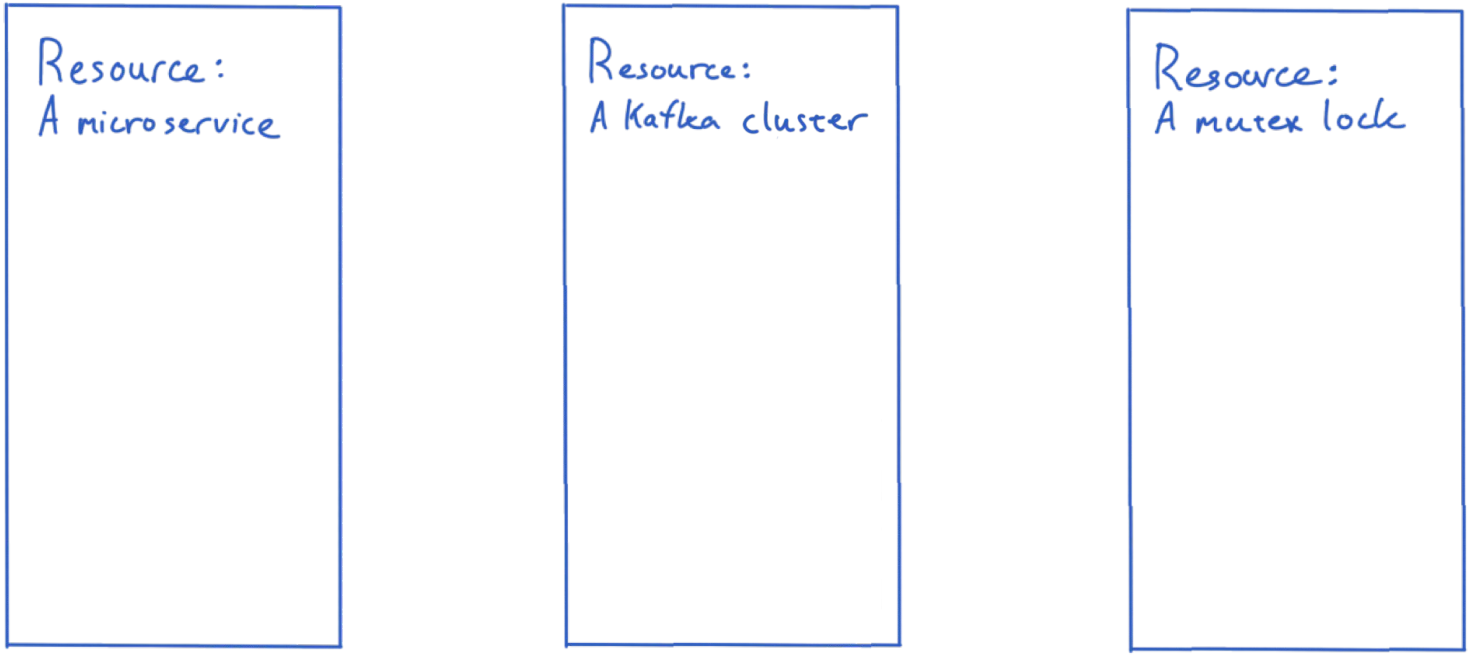

What are resources?

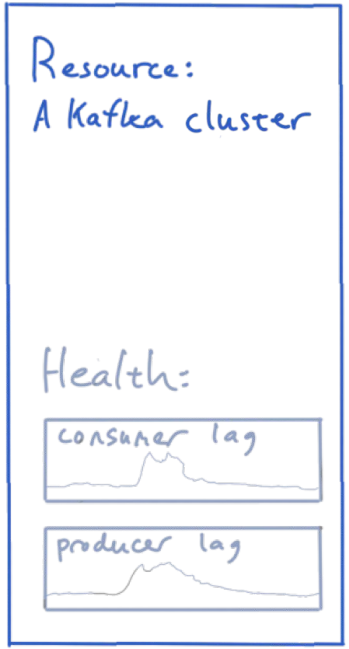

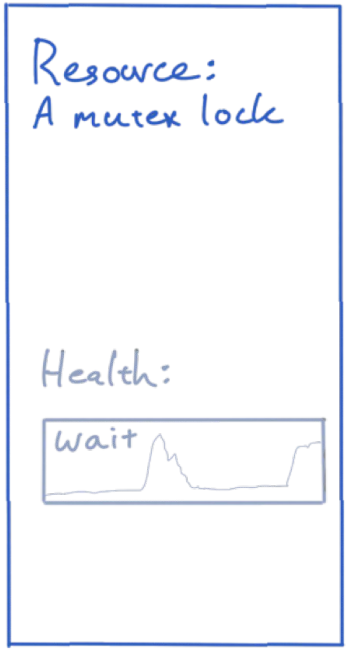

What are resources?

Resources are a hierarchy, too

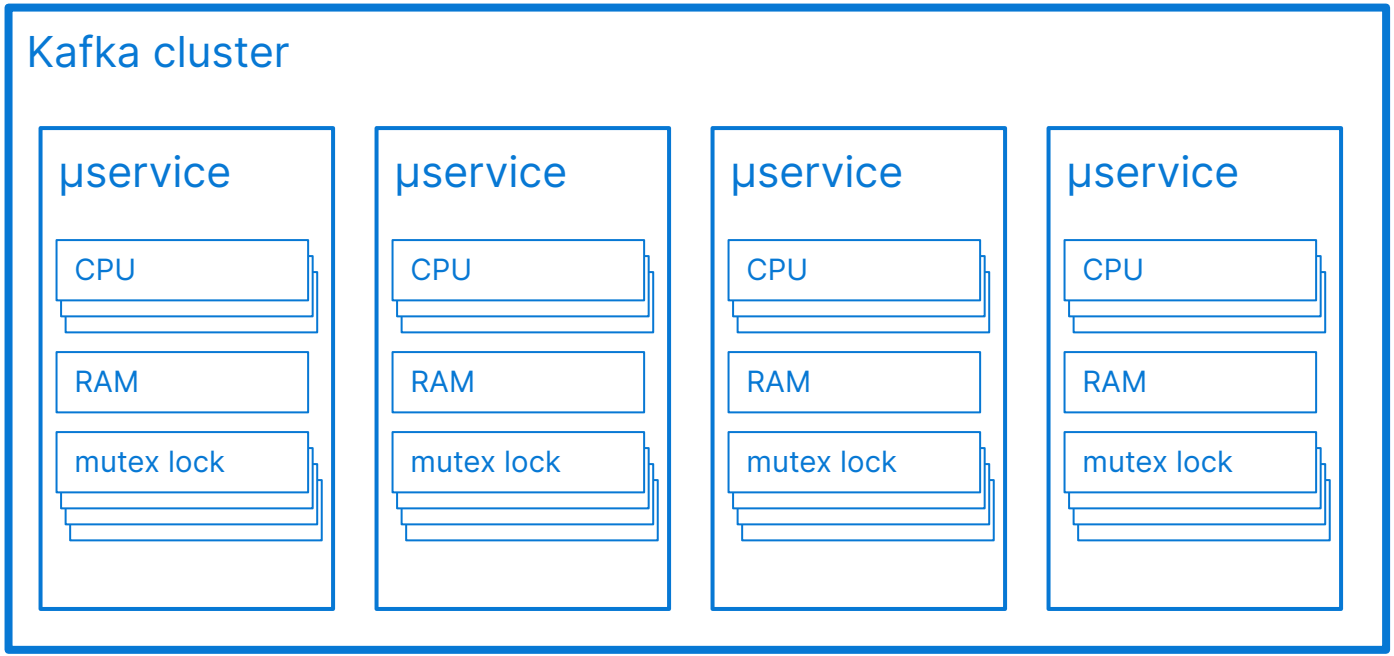

Part III

Co-dependency

Resources and transactions

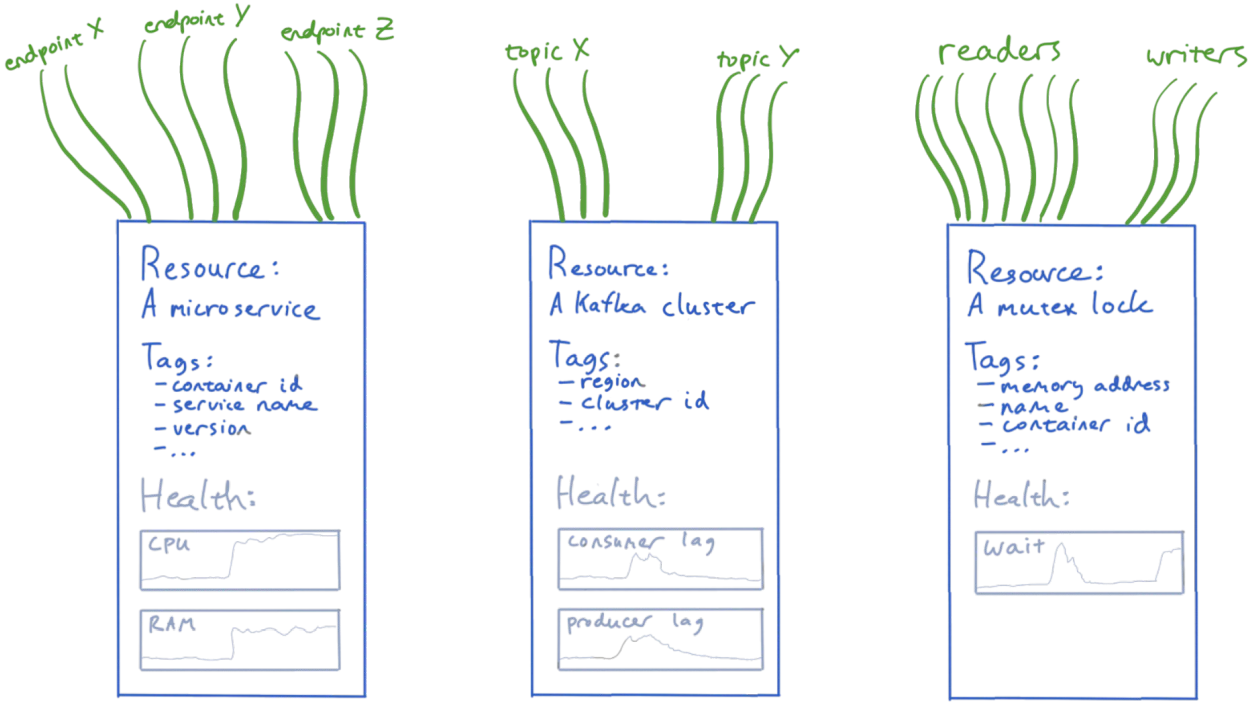

Resources and transactions

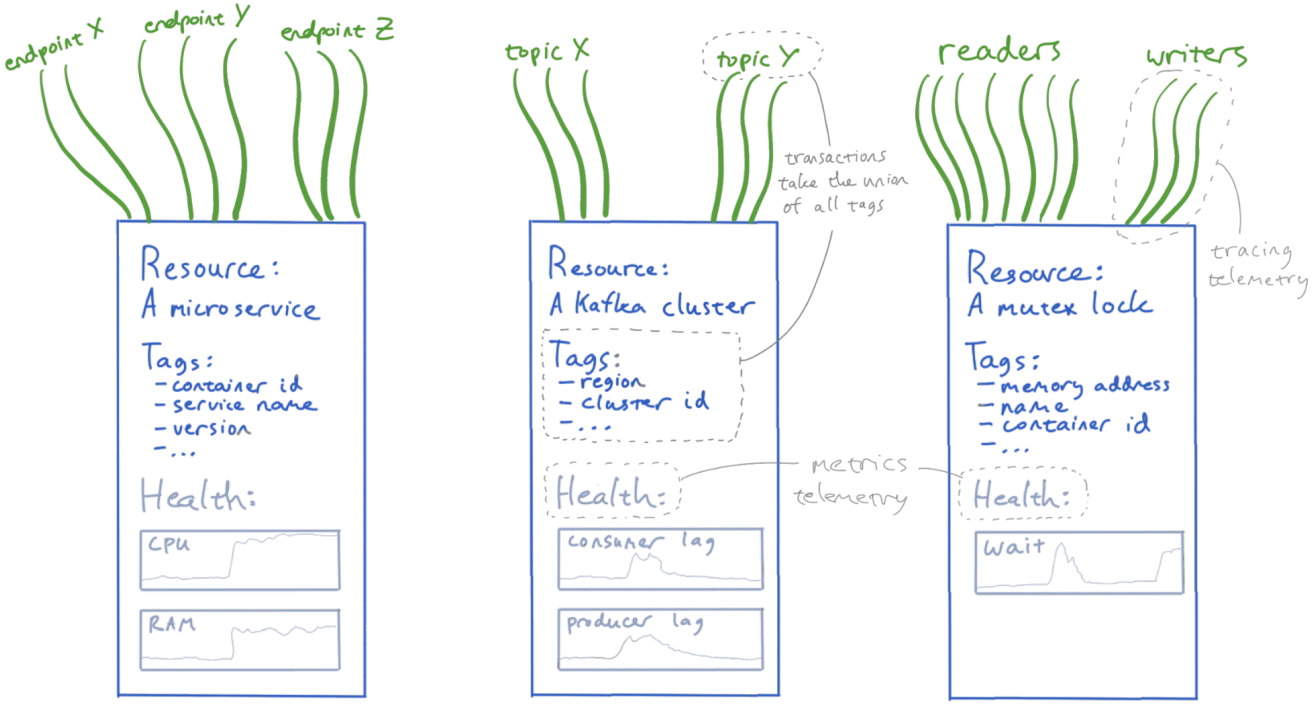

Resources and transactions are totally co-dependent

End-users only care about **Transactions!**

Operators only have control over **Resources!**

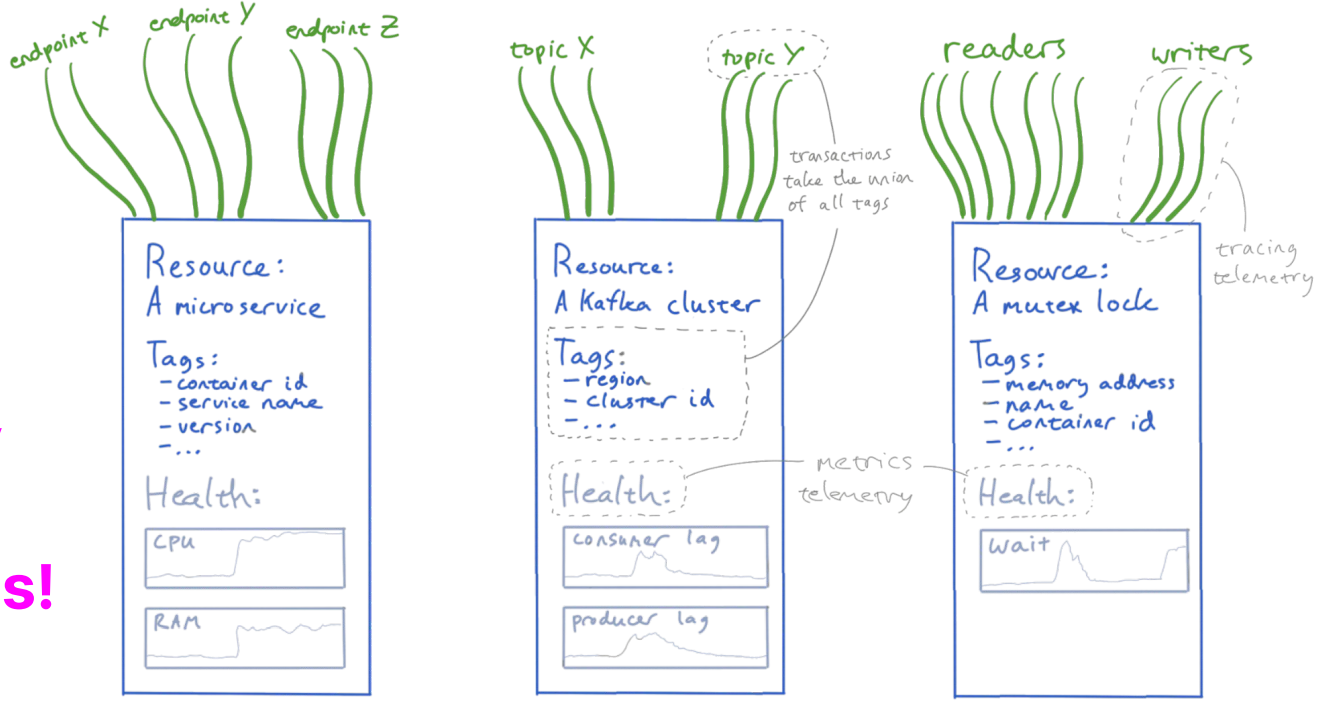

Resources and transactions are totally co-dependent

End-users and (dev)operators are totally co-dependent, too

End-users only care about **Transactions!**

Operators only have control over **Resources!**

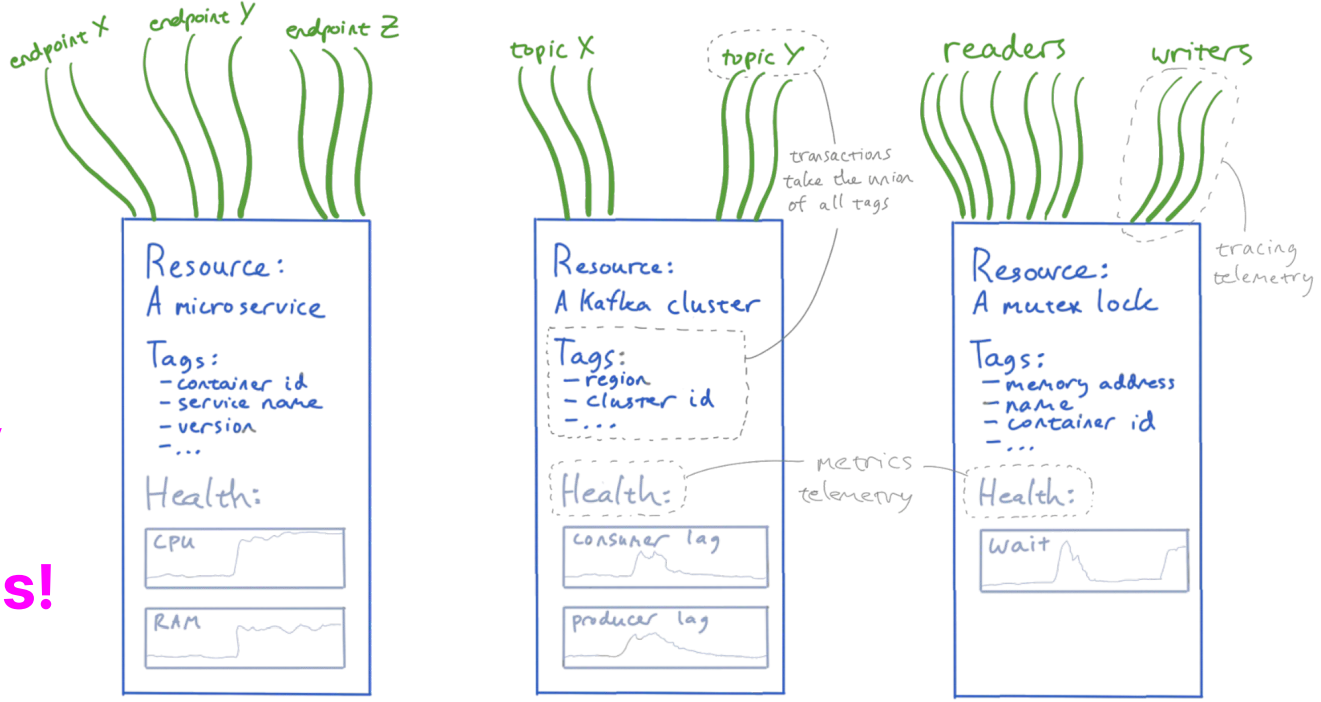

“What’s a DevOps eng / SRE to do???”

- Must pivot between resources and transactions
- … and thus across telemetry types (metrics ↔ traces)
- … in the aggregate
- … via tags
- … automatically
- … and without becoming an observability tooling expert

One helpful tool: **SLOs!**

- SLOs (“Service-Level Objectives”) are a hot topic these days 😎
- SLOs are goals ...
  - ... about **a set of Transactions** (e.g., error ratios, latency, etc)
  - ... scoped to **a set of Resources** (e.g., a microservice or DB)

In this way, **SLOs encompass Transactions, Resources, Operators, and End-Users.** No wonder they’re so helpful for observability!

Part IV

What does this look like in practice?


Feb 19 9:39 AM - 11:14 AM
Kafka consumer lag
Uh-oh...
10:30 AM
10:45 AM
11:00 AM

kafka.client.consumer.lag.messages

Cancel Save

Query

vanilla metrics query

Build a metrics query

kafka.client.consumer.lag.messages represent as a latest

Include: topic:public-usgs

Exclude: Search for tags

Group by: service and runinfo_host

Aggregate by the maximum value

+ Plot another metric + Add a formula

yikes... what changed??

View as Line chart

Feb 19 4:00 AM - 1:00 PM


Feb 19 4:00 AM - 1:00 PM
What caused this change?
Find possible causes through related traces

Have feedback? Share

kafka.client.consumer.lag.messages: What caused the change?

Selections for comparison

Display settings | Feb 19 4:00 AM - 1:00 PM

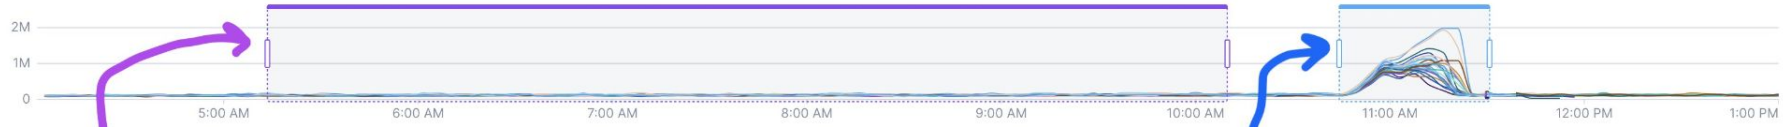

Baseline selected: Feb 19 5:13 AM - 10:10 AM
4 days ago
Collapse chart

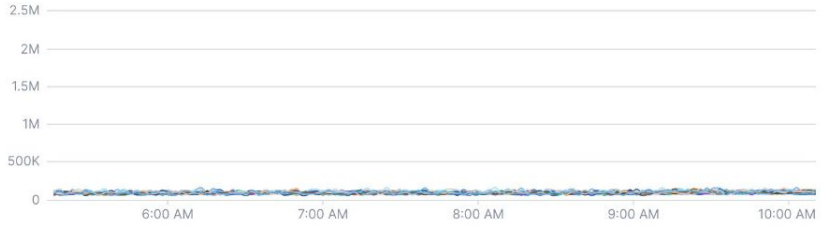

Deviation selected: Feb 19 10:44 AM - 11:31 AM
3 days ago
Collapse chart

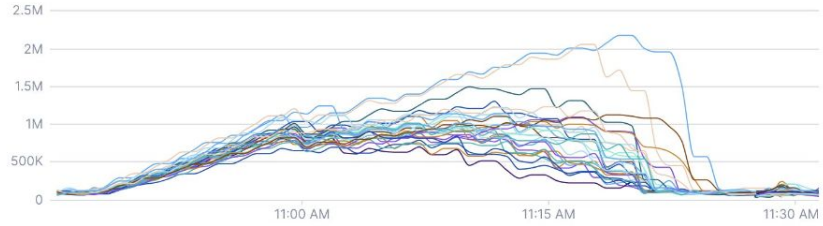

Biggest changes in aggie

Why aggie?

Sorted by Magnitude of change

aggie
aggregator.commit

aggie
aggregator.start_trace_assembly

Performance changes	Average	Average

Have feedback? Share

kafka.client.consumer.lag.messages: What caused the change?

Biggest changes in aggie

Why aggie?

Sorted by Magnitude of change

aggie
aggregator.commit

aggie
aggregator.start_trace_assembly

Performance changes

	Average	Average
Error	0%	0.93%
p95 Latency	124ms	1.42s
Rate	3.99K ops/s	4.05K ops/s

Most likely causes of performance changes

Baseline (990 traces) Deviation (760 traces)

Traces with project_fullname[redacted]	0.86% of traces	15.95% of traces
Traces with project_client_id:8c0072d91b0e97827988089104c60fdf	0.86% of traces	15.95% of traces
Traces with project_id:1753	0.86% of traces	15.95% of traces

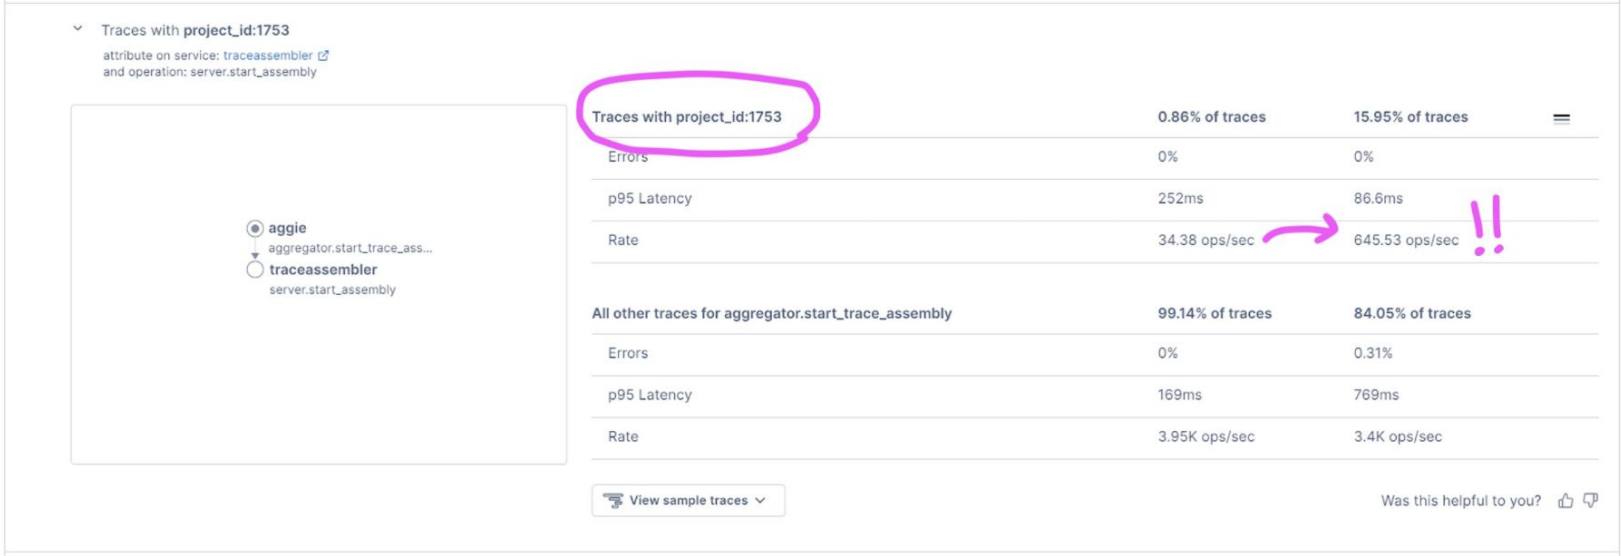
Traces with project_id:1753
attribute on service: traceassembler
and operation: server.start_assembly
aggie
aggregator.start_trace_ass...
traceassembler
server.start_assembly
Traces with project_id:1753
0.86% of traces
15.95% of traces
Errors
0%
0%
p95 Latency
252ms
86.6ms
Rate
34.38 ops/sec
645.53 ops/sec
!!
All other traces for aggregator.start_trace_assembly
99.14% of traces
84.05% of traces
Errors
0%
0.31%
p95 Latency
169ms
769ms
Rate
3.95K ops/sec
3.4K ops/sec
View sample traces
Was this helpful to you?

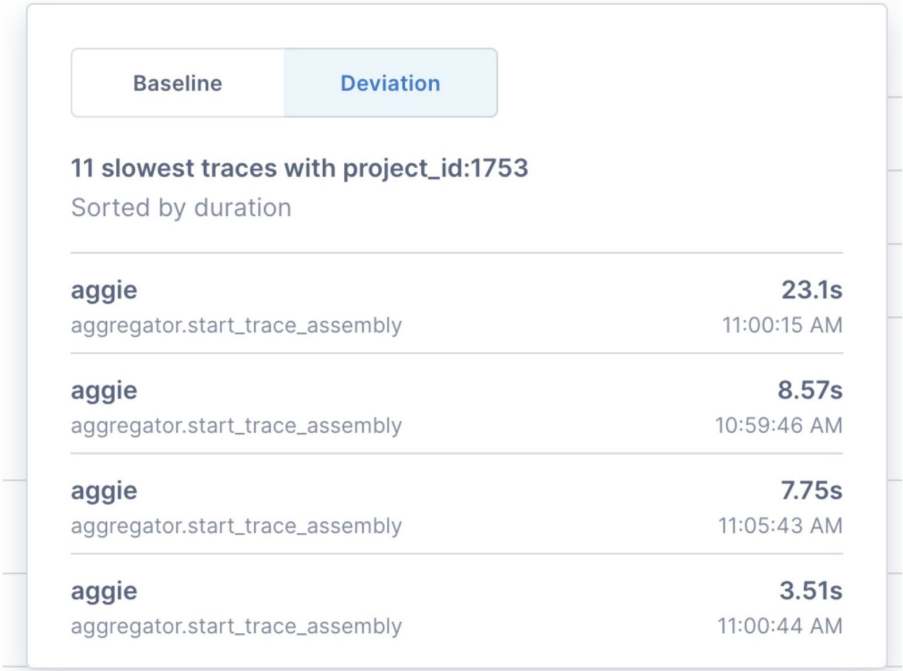
Baseline
Deviation
11 slowest traces with project_id:1753
Sorted by duration
aggie
23.1s
aggregator.start_trace_assembly
11:00:15 AM
aggie
8.57s
aggregator.start_trace_assembly
10:59:46 AM
aggie
7.75s
aggregator.start_trace_assembly
11:05:43 AM
aggie
3.51s
aggregator.start_trace_assembly
11:00:44 AM

Part V

Summing up...

- Transactions **traverse systems** and **use Resources**
- Users don’t give a s**t about your Resources
- DevOps can’t **do** s**t about individual Transactions
- We *must* be able to join Resources and Transactions to address the most important question in observability:

“What caused that change?”