

Differentiable Programming in Kotlin

Irene Dea

Differentiable Programming Team



Steffi Stumpos



Melissa Greuter



Samantha Andow



Irene Dea



Johann George



Shannon Yang



Alanna Tempest



Emilio Arroyo-Fang



Christy Warden



Neal Gafter



Xipeng Shen



Olin Shivers

What is Differentiable Programming?

Neural Networks are the beginning of a fundamental shift in how we write software.

- Andrej Karpathy, [Software 2.0](#)

*Differentiable Programming is a shift away from increasingly heavily parameterised [machine learning] models, to **simpler ones* that take more advantage of problem structure.***

- Mike Innes, [What Is Differentiable Programming](#)

Define networks procedurally in a data-dependent way (with loops and conditionals), allowing them to change dynamically as a function of the input data fed to them.

It's really very much like a regular program, except it's parameterized, automatically differentiated, and trainable/optimizable.

- Yann Lecun, [Deep Learning est mort. Vive Differentiable Programming!](#)

Motivations

- Popular frameworks are oriented towards traditional ML use cases
- There are many other use cases:
 - Computer graphics
 - Physics simulations
 - Probabilistic programming
 - And many more!

What We Need

- **Fast Language**
- **Automatic Differentiation (AD)**
- **Memory Safety**
- **Strong Types**
- **Static Compilation**

Performance
Usability
Flexibility

Our Approach

A compiler-aware framework for differentiation in Kotlin

- Customizable, extensible API
- Compile-time optimizations
- Compile-time shape checking

Why Kotlin?

- Usability
- Speed!
- Compiler Plugins
- Kotlin is *way* more than just Android
- Can target JVM, Native (LLVM), or Javascript

API: derivatives

Our Kotlin API supports scalar and tensor math, including derivatives. Given a Kotlin function, for example $f(x) = \sin(x)$

```
fun f(x: DValue) = sin(x)
```

To compute the first derivative $f'(x) = d/dx f(x)$ you write

```
fun fp(x: DValue) = derivative(x, ::f)
```

Our derivatives are computed at machine (float) precision, so $fp(x) == \cos(x)$.

```
fp(DFloat.PI) shouldBeExactly -1F
```

API: derivatives

We support both forward and reverse differentiation and, by nesting, higher-order derivatives.

```
fun fp1(x: DValue) = forwardDerivative(x, ::f)
```

```
fun fp2(x: DValue) = reverseDerivative(x, ::f)
```

```
fun fpp(x: DValue) = forwardDerivative(x,  
    { x -> reverseDerivative(x, ::f) })
```

API: more

We support many other components practically needed for AI applications:

- Tensors of arbitrary rank
- Sampling from random distributions
- Slicing, indexing, and concatenating tensors
- Computational layers commonly used for ML applications
 - Dense layer, convolution, norms, loss functions, ...
- Optimizers, Learning loops, ...

API: extensibility

Additionally, our API is designed to be customizable and extensible!

- User-defined differentiable types
- Trainable layers and components
- Optimizable by compiler plugins

API: performance

We are performant on a number of different use-cases.

API: performance

We are performant on a number of different use-cases.

- Kotlin
 - Fast custom logic, mixed workloads

API: performance

We are performant on a number of different use-cases.

- Kotlin
 - Fast custom logic, mixed workloads
- MKL-DNN
 - Traditional ML Models

API: performance

We are performant on a number of different use-cases.

- Kotlin
 - Fast custom logic, mixed workloads
- MKL-DNN
 - Traditional ML Models
- Sparse tensors
 - Graph problems, natural language processing, one hot or categorical data
 - Using Eigen library and a TACO-inspired data representation

API: performance

We are performant on a number of different use-cases.

- Kotlin
 - Fast custom logic, mixed workloads
- MKL-DNN
 - Traditional ML Models
- Sparse tensors
 - Graph problems, natural language processing, one hot or categorical data
 - Using Eigen library and a TACO-inspired data representation
- Compile-time optimizations
 - API is designed to be optimizable by a compiler plugin

AD Optimize Plugin

Problem

- AD compute tree is built and stored at runtime
- Scalar operations are boxed


Solution

- Inline differentiable computations
- Unbox scalars

AD Optimize Plugin: geometric series ($r = \frac{1}{2}$)

```
fun foo(a:DScalar): DScalar{  
    var y = DScalar(0f)  
    for (i in 0 until 1000) {  
        y += a / 2f.pow(i)  
    }  
    return y  
}  
  
val x:DScalar = DScalar(5f)  
val derivative =  
    reverseDerivative(x, ::foo)
```

- 1) Unbox Scalars
- 2) Inline derivative computation



AD Optimize Plugin: geometric series ($r = \frac{1}{2}$)

```
@ADOptimize
fun foo(a:DScalar): DScalar{
    var y = DScalar(0f)
    for (i in 0 until 1000) {
        y += a / 2f.pow(i)
    }
    return y
}
val x:DScalar = DScalar(5f)
val derivative =
    reverseDerivative(x, ::foo)
```

$$a / 2f^0 + a / 2f^1 + \dots$$

- 1) Unbox Scalars
- 2) Inline derivative computation

AD Optimize Plugin: geometric series ($r = \frac{1}{2}$)

```
@ADOptimize
fun foo(a:DScalar): DScalar{
    var y = DScalar(0f)
    for (i in 0 until 1000) {
        y += a / 2f.pow(i)
    }
    return y
}
val x:DScalar = DScalar(5f)
val derivative =
    reverseDerivative(x, ::foo)
```

$$a / 2f^0 + a / 2f^1 + \dots$$

But we can do more!

- 1) Unbox Scalars
- 2) Inline derivative computation

Coarsening Optimization: Concept

AD with Coarsening



**Algorithmic
Differentiation**

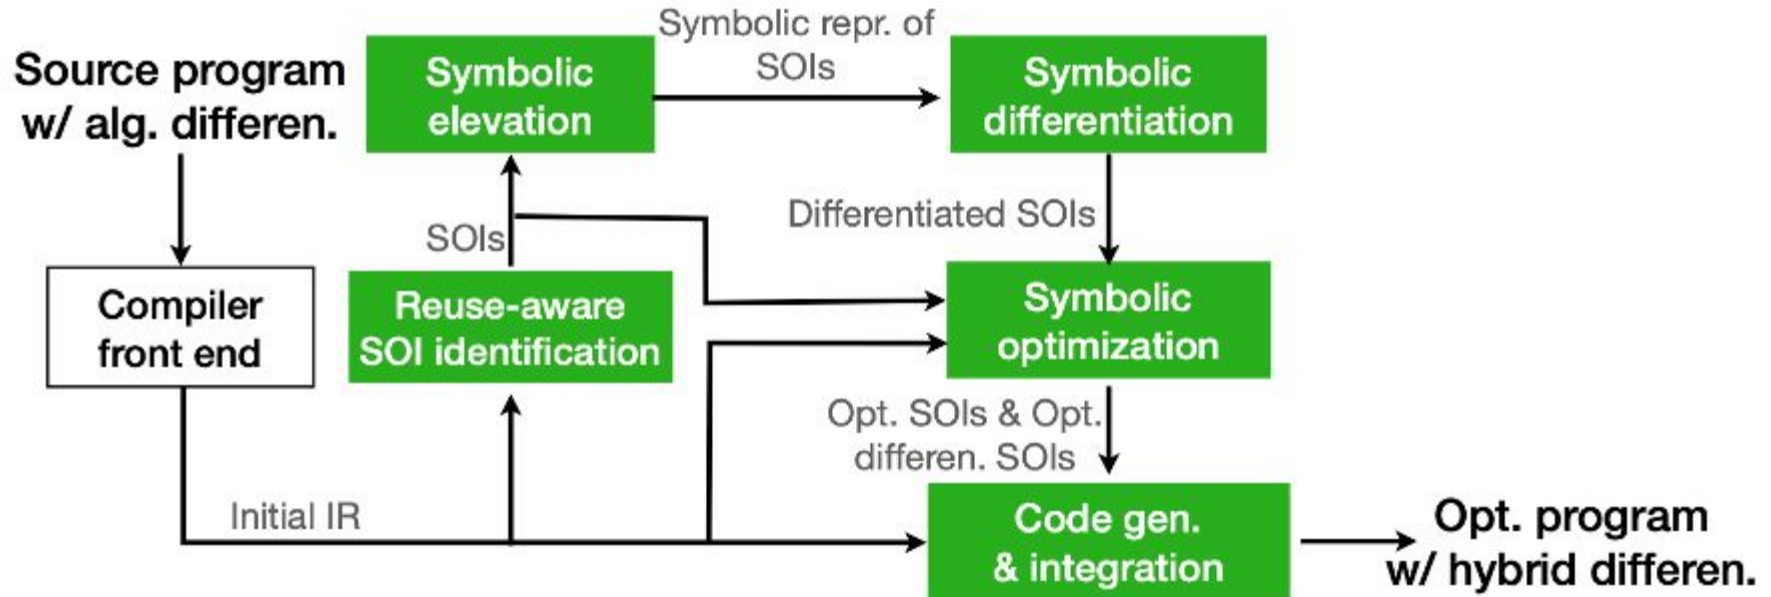
**Symbolic
Differentiation**

Best of both worlds

Finest granularity: Each operation

Largest granularity: Entire calculation

Coarsening Optimization: Workflow



Coarsening: geometric series ($r = \frac{1}{2}$)

```
fun foo(a:DScalar): DScalar{  
    var y = DScalar(0f)  
    for (i in 0 until 1000) {  
        y += a / 2f.pow(i)  
    }  
    return y  
}  
  
val x:DScalar = DScalar(5f)  
val derivative =  
    reverseDerivative(x, ::foo)
```

$$a * (1 - 0.5 ^ 1001) / (1 - 0.5)$$

For $r \neq 1$, the sum of the first $n+1$ terms of a geometric series, up to and including the r^n term, is

$$a + ar + ar^2 + ar^3 + \dots + ar^n = \sum_{k=0}^n ar^k = a \left(\frac{1 - r^{n+1}}{1 - r} \right),$$


Coarsening: geometric series ($r = \frac{1}{2}$)

```
fun foo(a:DScalar): DScalar{  
    return a * DScalar((1f - 0.5f.pow(1001) / (1 - 0.5f))  
}  
  
fun fooGrad(a:DScalar): DScalar{  
    return DScalar((1f - 0.5f.pow(1001) / (1 - 0.5f))  
}  
  
val x:DScalar = DScalar(5f)  
val derivative = fooGrad(x)
```

For $r \neq 1$, the sum of the first $n+1$ terms of a geometric series, up to and including the r^n term, is

$$a + ar + ar^2 + ar^3 + \dots + ar^n = \sum_{k=0}^n ar^k = a \left(\frac{1 - r^{n+1}}{1 - r} \right),$$


Time Reduction on Hookean Spring



	Primal (ms)	Gradients (ms)	Total (ms)
10 vertices	67=>0	93=>14	160=>15 (11X)
20 vertices	74=>0	106=>26	180=>27 (6.7X)
40 vertices	159=>0	221=>49	380=>51 (7.5X)

Time Reduction on Hookean Spring

Speedups of
1-2 orders of magnitude



	Primal (ms)	Gradients (ms)	Total (ms)
10 vertices	67=>0	93=>14	160=>15 (11X)
20 vertices	74=>0	106=>26	180=>27 (6.7X)
40 vertices	159=>0	221=>49	380=>51 (7.5X)

Static Shape Checking

- Compile-time tensor shape inference
- Compile-time tensor shape checking
- Real time feedback in IntelliJ
- Integration with our API
- User-defined shape functions

Static Shape Checking: example

```
@DeclareParams("A: _", "B: _", "C: _")
fun matmul(
    x: @ShapeOf("[A, B]") Tensor,
    y: @ShapeOf("[B, C]") Tensor
) : @ShapeOf("[A, C]") Tensor {
    ...
}
```

Static Shape Checking: example

```
@DeclareParams("A: _", "B: _", "C: _")
fun matmul(
    x: @ShapeOf("[A, B]") Tensor,
    y: @ShapeOf("[B, C]") Tensor
) : @ShapeOf("[A, C]") Tensor {
    ...
}

val a = Tensor(Shape(1, 2), ...) // [1,2]
val b = Tensor(Shape(2, 3), ...) // [2,3]
val res = matmul(a,b)           // [1,3]
```

Static Shape Checking: example

```
@DeclareParams("A: _", "B: _", "C: _")
fun matmul(
    x: @ShapeOf("[A, B]") Tensor,
    y: @ShapeOf("[B, C]") Tensor
) : @ShapeOf("[A, C]") Tensor {
    ...
}

val a = Tensor(Shape(1, 2), ...) // [1,2]
val b = Tensor(Shape(2, 3), ...) // [2,3]
val res = matmul(a,b)           // [1,3]

val badRes = matmul(b,a)      // ERROR: 2 != 3
```


Static Shape Checking: IntelliJ

```
val a = Tensor(Shape(...dims: 1,2), floatArrayOf(1f, 1f, 1f))  
val b = Tensor(Shape(...dims: 2,3), floatArrayOf(1f, 1f, 1f, 1f , 1f, 1f))
```

Tensor Shape([1, 3])

```
matmul(a,b)
```

```
val a = Tensor(Shape(...dims: 1,2), floatArrayOf(1f, 1f, 1f))  
val b = Tensor(Shape(...dims: 2,3), floatArrayOf(1f, 1f, 1f, 1f , 1f, 1f))
```

```
matmul(b,a)
```

[SHAPE_FUNCTION_ERROR] Shape Dimension Mismatch: 2 != 3

More Complex Shape Checking

`@ShapeFunction`

```
fun broadcast(a: Shape, b: Shape) : Shape? { ... }
```

`@DeclareParams("A: [____]", "B: [____]")`

```
fun add(a: @ShapeOf("A") Tensor, b: @ShapeOf("B") Tensor) :  
    @ShapeOf("broadcast(A, B)") Tensor { ... }
```

Use Case: Probabilistic Programming

- Collaboration with Facebook's PPL, Bean Machine
- Probabilistic Programming can benefit from
 - Higher order differentiation
 - Performant scalar support
 - Fast execution of native language
 - Sparse tensors
 - AD Optimize
 - Coarsening

Summary

- Performance
 - sparse tensors, MKL-DNN, Kotlin, optimization plugins
- Usability
 - functional API, static shape checking
- Flexibility
 - extensible API and plugins, probabilistic programming

Future Work

- Performance
 - Develop new optimizations enabled by compiler plugins
- Usability
 - Continue our work on static analyses
- Flexibility
 - Collaborate with users

Thank you!