
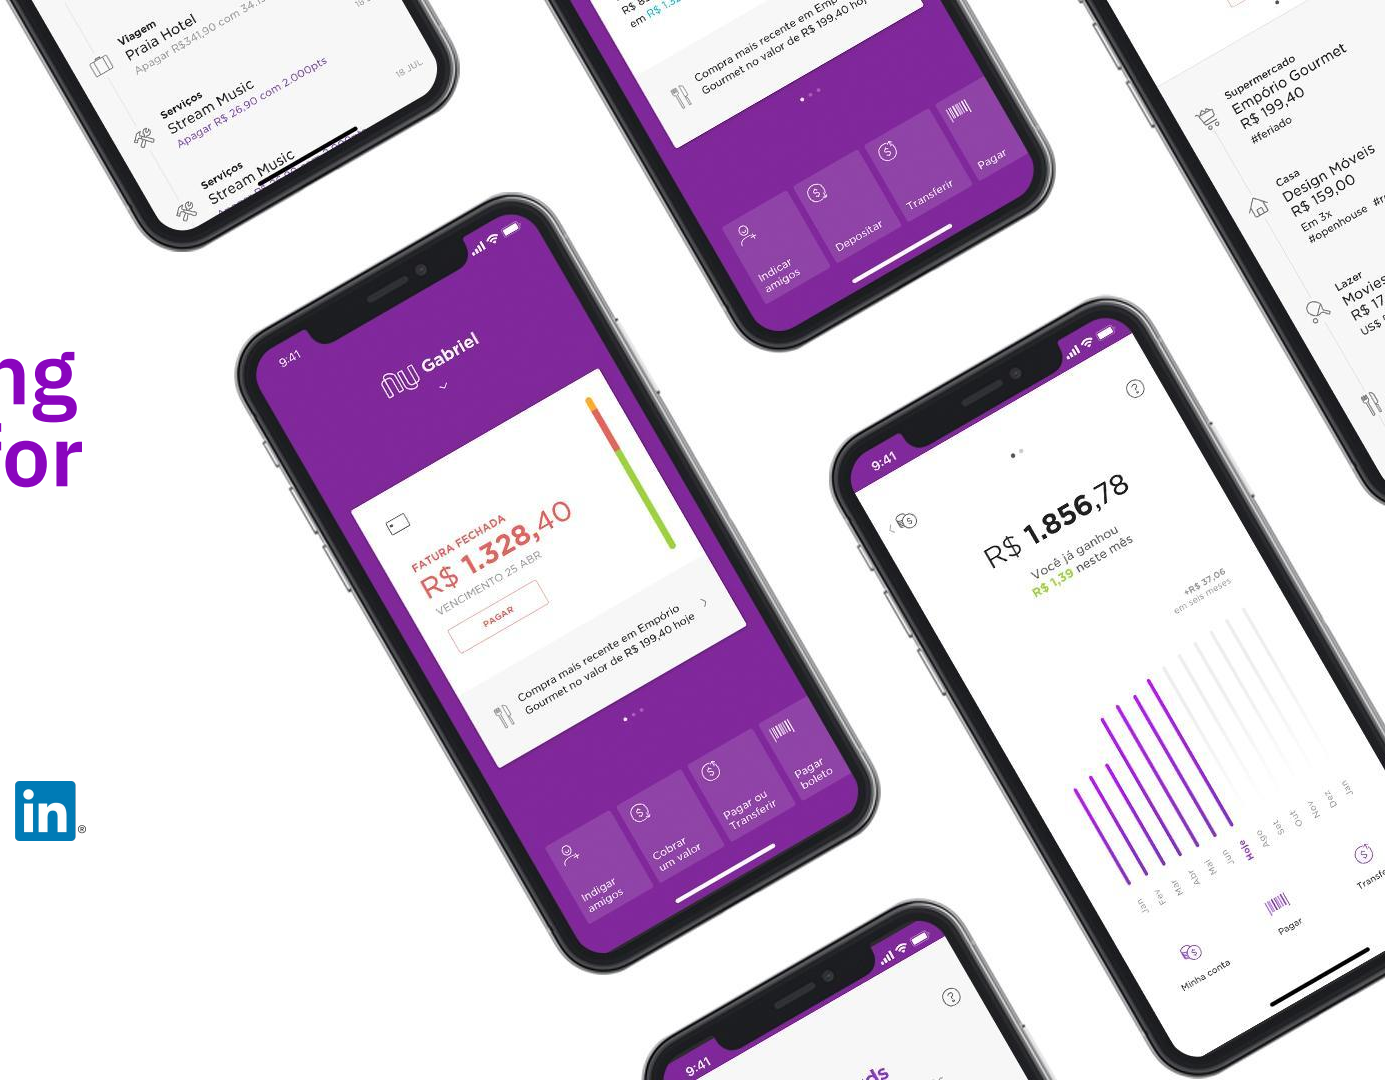


# Architecting Software for Leverage

Lucas Cavalcanti

@lucascs  

@lucascavalcantisantos 



# Leverage

*noun*

FINANCE

the ratio of a company's loan capital (**debt**) to the **value** of its common stock (equity).

*verb*

use borrowed capital for (an **investment**), expecting the profits made to be greater than the interest payable.



ARCHITECTING FOR LEVERAGE

# What to take away from this talk

Example of architectural decisions taken during different stages of Nubank's trajectory, aiming at the highest leverage aspect at the time  
You may be in a similar position on your current company

# Who am I?




Lucas Cavalcanti  
Principal Software Engineer  
@ Nubank since 2013  
Based in São Paulo - Brazil

Nubank is the **leading fintech in Latin America**,  
born to eliminate complexity and empower  
customers to take control of their money.



# Growing rapidly in a complex domain



**35M**

Customers

**1B**

HTTP Requests/day

**1B**

Kafka Messages/day

**10s**

Deploys/day

**700**

Microservices

**16**

Shards

**700**

Engineers

**130**

teams

# Agenda

Startup time: Time to market and Feedback

Growth time: Resilience and adaptability

Consolidation time: Reliability and observability


Expansion time: Flexibility and extensibility

# Startup time



## STARTUP TIME

Late 2013 to  
early 2015



## GREENFIELD PROJECT

That magical moment when you get to choose any technology

## FIRST PRODUCT

A digital first credit card with no fees and real time experience on the app

## LIMITED RESOURCES

Just few people and a limited cash source to burn

## LICENSE DEADLINE

If not operating by May 2014, we'd have to apply for a license with would take up to 2 years to be granted

## SMALL "OFFICE"

A little house in a quiet neighborhood in São Paulo

## UNKNOWNNS

No previous domain knowledge. Uncertain if product would get traction

# Technology Choices

VALUE: Time to market

LEVERAGE TYPE: Maximizing work not to be done; Containing complexity

✓ DATABASE: DATOMIC

- Immutable ledger database of facts
- Auditing for free
- Updates preserve history
- Querying database at any point in time

✓ LANGUAGE: CLOJURE

- Runs on the JVM: leverage the whole Java ecosystem
- Immutability by default
- Simple made easy
- Functional Programming close to finance
- Hexagonal architecture

✓ MESSAGING: KAFKA

- Persistent log of messages (with a TTL)
- Ability for resetting offsets so we can reprocess old messages
- Strong durability guarantees
- Topics partitioned by default

DEBT: Niche and unconsolidated technologies.  
Hard to find people with previous experience

# Vendors

**VALUE:** Time to market

**LEVERAGE TYPE:** Maximizing work not to be done; Build vs Buy

✓ **CLOUD: AWS**

CloudFormation for deploy automation  
DynamoDB for scalable storage for Datomic  
Easy to scale services

✓ **CREDIT CARD PROCESSOR**

Off the shelf solution to run a custom Credit Card  
2 months to integrate  
MasterCard setup and licensing already done

**DEBT:** Limited by the processor's ability to scale and respond to problems

# Practices

**VALUE:** Fast and early feedback

**LEVERAGE TYPE:** foundation to build on top of it faster

✓ **CI/CD**

- Github + Pull Requests
- Unit and integration tests per service
- E2E tests with all services
- Baked images (pre-docker)
- Test, Staging and Production envs

✓ **FAULT TOLERANCE**

- Rudimentary transient monitoring
- Service's health
- Immutable infrastructure

✓ **MICROSERVICES**


- Domain is intrinsically complex
- Breaking complexity in smaller services would help contain it

**INVESTMENT:** Time to build the foundation, which is not always available in startup phases

# Growth time

GROWTH TIME

# 2015 and 2016



## FAST GROWTH

Customer base grew way faster than any projection

## VENDOR NOT SCALING

Credit card processor not responding fast enough to the increase in number of customers

## "OFFICE" NOT SCALING

Moved to a small 3-story office building near Ibirapuera Park in São Paulo which would fit 150 people

## TECH NOT SCALING

Scaling horizontally would only get us so far. Bottlenecks started appearing

# Practices

**VALUES:** Scalability, Fault tolerance

**LEVERAGE TYPE:** Avoiding/delaying optimization

## ✓ INFRASTRUCTURE SHARDING

- Sharding the whole infrastructure
- Scalability units
- Limited blast radius on failures
- If shards are small enough, no need for optimization

## ✓ MORE CI/CD

- Frequent automatic deploys
- E2e taking up to 1h to run
- Replaced e2e with consumer-driven contract tests
- Docker images rather than EC2 images

**INVESTMENT:** year long project to implement and rollout sharding.  
Design and creation of a whole new tool

**DEBT:** project took longer than expected and first shard was bigger than expected and was a special shard for a long time  
Each shard has a minimal AWS cost, regardless of number of customers

# In-housing

**VALUES:** Resilience, Scalability, Flexibility

**LEVERAGE TYPE:** Owning your own destiny

## ✓ PROCESSING IN-HOUSING

Processing CC transactions is the core of the business  
Vendor didn't scale as expected  
Any change would take months of back and forth to implement

## ✓ CUSTOMER SUPPORT IN-HOUSING

Providing the best customer support in the industry was key to success  
Existing solutions didn't have enough flexibility or were too hard to configure

**INVESTMENT:** 18mo long project to bring each feature in-house and migrate from previous system


**DEBT:** Long period without any major product changes



# Consolidation time

CONSOLIDATION TIME

# 2017 and 2018



## SCALABLE YET UNSTABLE

Although sharding helped a lot on the scaling part, we got to the point where every little corner case would happen, several times

## SECOND PRODUCT

With Credit Card consolidating, we moved to a checking account product

## OFFICE NOT SCALING, AGAIN

Moved to a 8-story office building near Av. Paulista in São Paulo, which would fit over 1000 people

## BIG DATA

Aggregating data from all services and shards became of a huge importance

# Technology

**VALUES:** Scalability, Adaptability, Observability

**LEVERAGE TYPE:** Ease of infrastructure changes. Cloud tools

## ✓ KUBERNETES

Ecosystem of infrastructure tools  
With high number of services, it  
scales better than AWS  
CloudFormation

## PROMETHEUS + GRAFANA

Collecting real time metrics  
Operational dashboards  
Inflow of metrics for other tools,  
like OpsGenie, Slack, CI/CD  
canary deploys

**INVESTMENT:** year long project to setup and migrate shard by shard to k8s

**DEBT:** constantly hitting AWS limits and spending \$\$\$ with duplicated infrastructure until project was done

# Internal Tools

**VALUES:** Resilience, Observability

**LEVERAGE TYPE:** Engineering operational productivity

✓ **NuCLI**

Repository of command line tools  
evolved by all engineers  
Most common operations (like  
restarting a service or curl'ing with  
proper credentials) easy at hand

✓ **DECLARATIVE INFRA**

Repository of declarative  
resources configuration  
Automatic update as soon as  
change hits the main branch at  
github

**INVESTMENT:** dedicated team to curate, maintain and ensure that  
changes get applied properly

# Data

**VALUES:** Observability, Consistency

**LEVERAGE TYPE:** Support for pretty much every decision in the company

## ✓ SCALA + SPARK

Scalable tools to extract data from all service's databases from all the shards, and then compose them into higher level business definitions

## ✓ ETL

Repository of dataset definitions  
Contributions from people of diverse functions in the company  
Outputs to data warehouses  
Integrated with BI tools  
Support for ML

## CONSISTENCY CHECKS

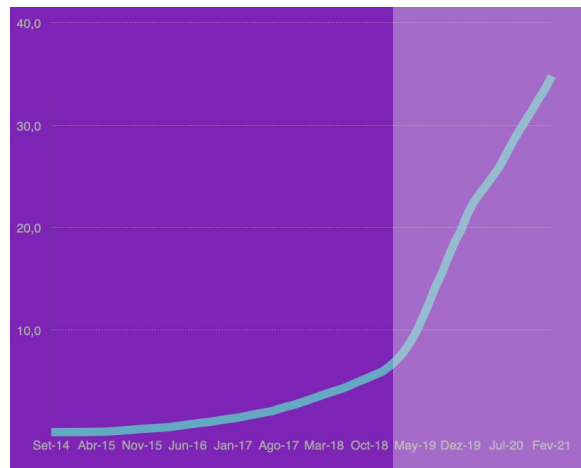
Automated controls that check for inconsistent data and alarm above a threshold  
Tool for finding partially completed or failed distributed transactions

**INVESTMENT:** year long project to create initial versions.  
Dedicated team maintain and keep ETL running in a reasonable time  
\$\$\$ on AWS to run the ETL process every day

# Expansion time

EXPANSION TIME

# 2019 to present



## PRODUCTS FOR EVERYONE

A product offering for everyone that applied

## MANY COUNTRIES

Launched products in MX and CO

## MANY OFFICES

Offices in BR, DE, MX, CO, AR and US, over 3000 employees

## MANY PRODUCTS

Product portfolio growing. Several configurations of existing products

## ACQUISITIONS

Acquired Plataformatec, Cognitect and Easyinvest

# Horizontal Platforms

**VALUES:** Extensibility, Productivity

**LEVERAGE TYPE:** Technology specialist teams building abstracted tools for the others to use

## ✓ MOBILE + WEB

Flutter + Dart for mobile  
ClojureScript + Fulcro for web  
NuDS + component libraries  
BFFs with GraphQL or REST

## ✓ INFRASTRUCTURE

Mostly declarative definitions  
NuCLI for common operations  
All horizontal aspects solved centrally, e.g, fault tolerance, auto scaling, deployment, monitoring

## EXPERIMENTATION

Infrastructure to run experiments with feature toggles, A/B testing and monitoring of KPIs

**INVESTMENT:** dedicated teams creating and maintaining component libraries, tools, documentation and guidelines, with specialists in those technologies so no others need to be specialist to be able to extend and create new things on them



# Business Platforms

**VALUES:** Extensibility, Productivity, Containing complexity

**LEVERAGE TYPE:** Domain specialist teams building abstracted APIs for the others to use.  
Endless possibilities of products built on top of the platforms



## BANKING AS A SERVICE

Credit platform to abstract loan issuing, interest generation, reporting and accounting  
Future Assets and Payments/Transfers platforms  
OpenBanking



## INTERNATIONAL CREDIT CARD

Credit Card system broken into platforms for most relevant subdomains: Limits, Billing, Transaction processing, Debts, and others

## ACQUISITION

Platform for offering products to customers  
Flexible acquisition processes

**INVESTMENT:** Deep and long design sessions with domain experts to build the long term vision for each business platforms and effort to migrate from legacy systems to the platforms

# Recap

**Startup time:** Delay writing code and build foundation

**Growth time:** In-housing and sharding

**Consolidation time:** Infrastructure and data

**Expansion time:** Horizontal and business platforms

# Thanks :)

Lucas Cavalcanti

@lucascs  

@lucascavalcantisantos 

