

# Unified MLOps: Feature Stores and Model Deployment

---


Monte Zweben  
Co-Founder & CEO

QCon 2021




# Why Do You Need a Feature Store?

1-5 models deployed



100's models deployed

# Why Is This So Hard?



## Productivity

Feature engineering consumes everyone



## Predictive Accuracy

Training and serving pipelines are often inconsistent and wrong



## Model Governance

Why did the model do that?


# How do you use a Feature Store?

Feature Reuse

Search...




Create Training Sets



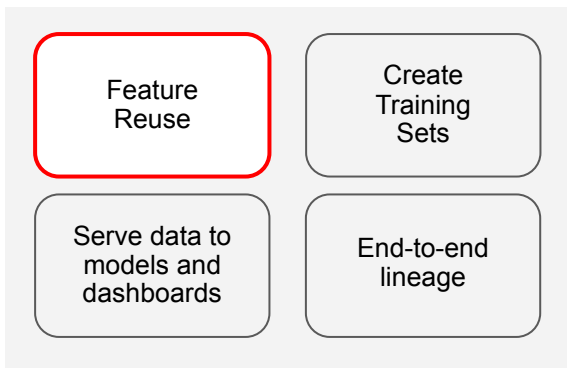
Serve data to models and  
dashboards



End-to-end lineage



# Feature Reuse



## Without a Feature Store

Data Scientists define features in Python, engineers then translate into robust SQL or Spark pipelines

Duplicate code and infrastructure for redundant feature engineering pipelines

## With a Feature Store

Data Scientists easily define production ready features


Single line of Feature Store API to reuse features

# Feature Search

Index	NAME	FEATURE_TYPE	DESCRIPTION	SCHEMA_NAME	TABLE_NAME	SETDESCRIPTION
0	CUSTOMER_RFM_DELICATESSEN_RATE_1W	C	Last weeks units purchased count in the Deli category.	RETAIL_FS	CUSTOMER_RFM_BY_CATEGORY	Describes customer by aggregating their purchases by category over multiple time window
1	CUSTOMER_RFM_DELICATESSEN_RATE_2W	C	Last 2 weeks units purchased count in the Deli category.	RETAIL_FS	CUSTOMER_RFM_BY_CATEGORY	Describes customer by aggregating their purchases by category over multiple time window
2	CUSTOMER_RFM_DELICATESSEN_RATE_4W	C	Last 4 weeks units purchased count in the Deli category.	RETAIL_FS	CUSTOMER_RFM_BY_CATEGORY	Describes customer by aggregating their purchases by category over multiple time window
3	CUSTOMER_RFM_DELICATESSEN_RATE_8W	C	Last 8 weeks units purchased count in the Deli category.	RETAIL_FS	CUSTOMER_RFM_BY_CATEGORY	Describes customer by aggregating their purchases by category over multiple time window
4	CUSTOMER_RFM_DELICATESSEN_RATE_16W	C	Last 16 weeks units purchased count in the Deli category.	RETAIL_FS	CUSTOMER_RFM_BY_CATEGORY	Describes customer by aggregating their purchases by category over multiple time window
5	CUSTOMER_RFM_DELICATESSEN_RATE_32W	C	Last 32 weeks units purchased count in the Deli category.	RETAIL_FS	CUSTOMER_RFM_BY_CATEGORY	Describes customer by aggregating their purchases by category over multiple time window

```
requiring feature data...
mean: 1.62    std: 19.57
median: 0.0
range: -38.0 to 1080.0
```

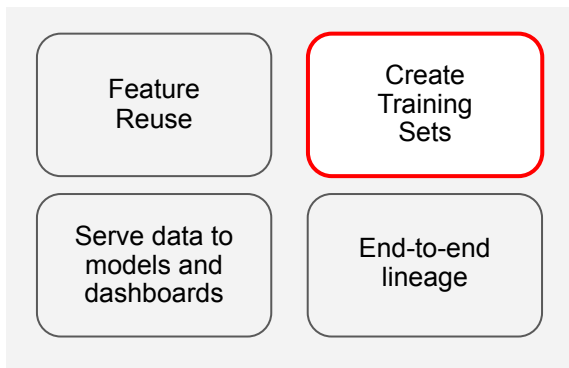
**Feature - RETAIL\_FS.CUSTOMER\_RFM\_BY\_CATEGORY -  
CUSTOMER\_RFM\_DELICATESSEN\_RATE\_1W Distribution**



# Automatically Generate Common Feature Transformations

```
39 #####
40
41 start_time = '2020-12-28T00:00:00'
42 schedule_interval = AggWindow.get_window(7,AggWindow.DAY)
43
44
45 backfill_start = datetime.strptime('2019-04-01 00:00:00', '%Y-%m-%d %H:%M:%S')
46 backfill_interval = schedule_interval
47 |
48 #####
49
50 fs.create_aggregation_feature_set_from_source(
51     source_name, 'ecommerce_fs', 'customer_purchases', start_time=start_time,
52     schedule_interval=schedule_interval, backfill_start_time=backfill_start,
53     backfill_interval=backfill_interval,
54     aggregations = [
55         FeatureAggregation(feature_name_prefix = 'PRODUCE_REVENUE',column_name = 'PRODUCE_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
56         FeatureAggregation(feature_name_prefix = 'DELI_REVENUE',column_name = 'DELI_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
57         FeatureAggregation(feature_name_prefix = 'DAIRY_REVENUE',column_name = 'DAIRY_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
58
59         FeatureAggregation(feature_name_prefix = 'CARROTS_REVENUE',column_name = 'CARROTS_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
60         FeatureAggregation(feature_name_prefix = 'APPLES_REVENUE',column_name = 'APPLES_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
61         FeatureAggregation(feature_name_prefix = 'BANANAS_REVENUE',column_name = 'BANANAS_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
62         FeatureAggregation(feature_name_prefix = 'GRAPES_REVENUE',column_name = 'GRAPES_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
63         FeatureAggregation(feature_name_prefix = 'TURKEY_REVENUE',column_name = 'TURKEY_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
64         FeatureAggregation(feature_name_prefix = 'CHICKEN_REVENUE',column_name = 'CHICKEN_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
65         FeatureAggregation(feature_name_prefix = 'BEEF_REVENUE',column_name = 'BEEF_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
66         FeatureAggregation(feature_name_prefix = 'SKIM_MILK_REVENUE',column_name = 'SKIM_MILK_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
67         FeatureAggregation(feature_name_prefix = 'SOY_MILK_REVENUE',column_name = 'SOY_MILK_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0),
68         FeatureAggregation(feature_name_prefix = 'CHEESE_REVENUE',column_name = 'CHEESE_REVENUE',agg_functions=['sum'], agg_windows=['1d','7d', '30d', '365d'],agg_default_value=0)
69     ]
70 )
71
```

# Create Training Set



## Without a Feature Store

Hundreds of lines of complex and error prone SQL to join features and labels correctly

Feature values of the past may be lost, making it nearly impossible to build training sets in the future

## With a Feature Store

Simply specify a training label and join keys to automatically create training sets

Feature values are automatically versioned for use in future models



# Training Set Creation

```
1 sql = """
2 SELECT ltv.CUSTOMERID,
3        ((w.WEEK_END_DATE - ltv.CUSTOMER_START_DATE)/ 7) CUSTOMERWEEK,
4        CAST(w.WEEK_END_DATE as TIMESTAMP) CUSTOMER_TS,
5        ltv.CUSTOMER_LIFETIME_VALUE as CUSTOMER_LTV
6 FROM retail_rfm.weeks w --splice-properties useSpark=True
7 INNER JOIN
8     retail_fs.customer_lifetime ltv
9     ON w.WEEK_END_DATE > ltv.CUSTOMER_START_DATE AND w.WEEK_END_DATE <= ltv.CUSTOMER_START_DATE + 28 --only first 4 weeks
10 """
11
12 pks = ['CUSTOMERID', 'CUSTOMERWEEK'] # Each unique training row is identified by the customer and their week of spending activity
13 join_keys = ['CUSTOMERID'] # This is the primary key of the Feature Sets that we want to join to
14
15 fs.create_training_view(
16     'Customer_Lifetime_Value',
17     sql=sql,
18     primary_keys=pks,
19     join_keys=join_keys,
20     ts_col = 'CUSTOMER_TS', # How we join each unique row with our eventual Features
21     label_col='CUSTOMER_LTV', # The thing we want to predict
22     desc = 'The current (as of queried) lifetime value of each customer per week of being a customer'
23 )
```

Registering Training View Customer\_Lifetime\_Value in the Feature Store

## Easily extract all features

*Every time this code is re-run you have access to the most up-to-date features*

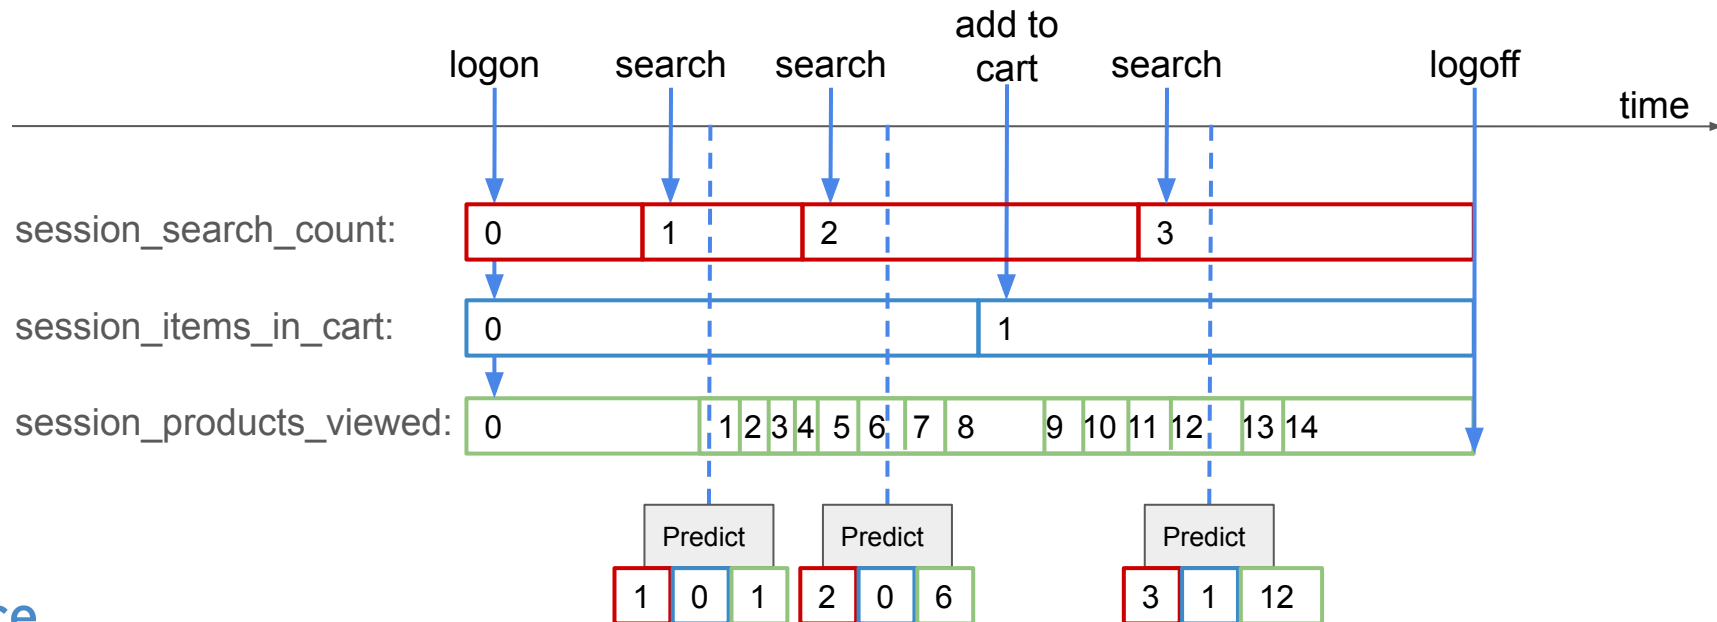
```
1 #Spark Dataframe
2 all_features = fs.get_training_set_from_view('Customer_Lifetime_Value')
3 all_features.limit(8).toPandas()
```

CUSTOMER_LIFETIME_ITEMS_PER_ACTIVE_DAY	CUSTOMER_LIFETIME_REVENUE_PER_ACTIVE_DAY	CUSTOMER_LIFETIME_DAYS	CUSTOMER_DAYS_SINCE_PURCHASE	CUSTOMER_LIFETIME_VALUE	CUSTOMER_START_DATE	CUSTOMER_L
137.0	219.350	522	522	219.35	2019-11-04	219.35
137.0	219.350	522	522	219.35	2019-11-04	219.35
137.0	219.350	522	522	219.35	2019-11-04	219.35
137.0	219.350	522	522	219.35	2019-11-04	219.35

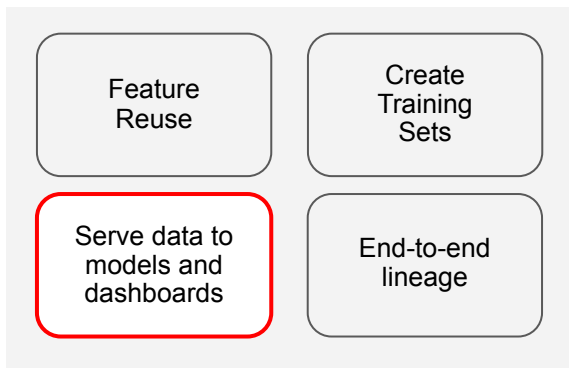
# Point-in-Time Correctness

Training should be done with the features available at model run time.

**The problem:** features change over time.



# Serving Features



## Without a Feature Store

Create bespoke pipelines that feed a key-value store

Inevitable inconsistency between two separate databases means the features you need aren't always available or consistent with training data

## With a Feature Store

Serve the same features used to train models to deployed models

ACID compliant triggers ensures that data used for training is always available for serving

# Feature Serving

```
1 feature_vector_sql = fs.get_feature_vector(features=features_list, return_sql=True, join_key_values={'customerid':'14235'})
2 print(feature_vector_sql)
```

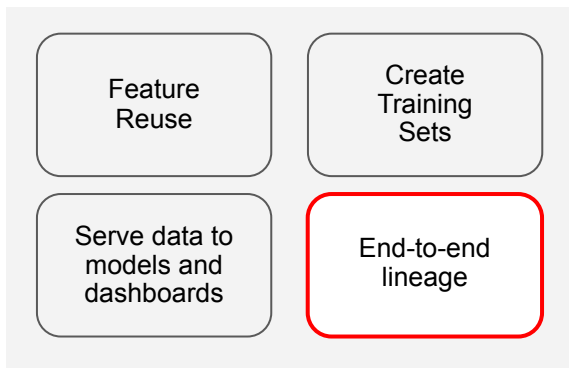
```
SELECT fset2.customerid, fset2.customer_rfm_home_revn_rate_4w, fset2.customer_rfm_toys_rate_2w, fset2.customer_rfm_jewelry_rate_1w, fset2.customer_rfm_total_revn_rate_4w, fset2.customer_rfm_delicatessen_r
ate_4w, fset2.customer_rfm_home_revn_rate_2w, fset2.customer_rfm_school_supplies_rate_2w, fset3.customer_lifetime_days, fset2.customer_rfm_kitchen_revn_rate_4w, fset2.customer_rfm_total_rate_4w
FROM retail_fs.customer_rfm_by_category AS fset2, retail_fs.customer_lifetime AS fset3
WHERE fset2.customerid = '14235' AND fset3.customerid = '14235'
```

```
1 --time
2 --sql
3 SELECT CUSTOMER_RFM_HOME_DECOR_RATE_4W,CUSTOMER_RFM_KITCHEN_RATE_1W,CUSTOMER_RFM_TOTAL_RATE_4W,CUSTOMER_RFM_DELICATESSEN_REVN_RATE_2W,CUSTOMER_RFM_HOME_REVN_RATE_1W,CUSTOMER_RFM_HOME_REVN_RATE_4W,CUSTOM
4 FROM retail_fs.CUSTOMER_RFM_BY_CATEGORY fset1,
5     retail_fs.CUSTOMER_LIFETIME fset10
6 WHERE fset1.CUSTOMERID = 14235 AND fset10.CUSTOMERID = 14235
```

i index	Key	Value
0	CUSTOMER_RFM_HOME_DECOR_RATE_4W	0
1	CUSTOMER_RFM_KITCHEN_RATE_1W	0
2	CUSTOMER_RFM_TOTAL_RATE_4W	0
3	CUSTOMER_RFM_DELICATESSEN_REVN_RATE_2W	0
4	CUSTOMER_RFM_HOME_REVN_RATE_1W	0
5	CUSTOMER_RFM_HOME_REVN_RATE_4W	0
6	CUSTOMER_RFM_HOME_DECOR_REVN_RATE_4W	0
7	CUSTOMER_RFM_KITCHEN_REVN_RATE_4W	0
8	CUSTOMER_RFM_TOTAL_REVN_RATE_4W	0
9	CUSTOMER_LIFETIME_DAYS	612

CPU times: user 6.69 ms, sys: 1.13 ms, total: 7.82 ms

# End-to-End Governance



## Without a Feature Store

Impossible or prohibitively expensive to find the training set used for a model


Search through API logs to determine what features were used and what predictions were made

## With a Feature Store

Instantly return the exact training set used for a model without having to persist the dataset


Identify the features served to deployed models easily with a built in evaluation store

# End-to-End Governance



CHOOSE WORKSPACE

Erin Space ACTIVE

 **Erin Splice**  
Jack Ploshnick

QUICK START

MACHINE LEARNING

Notebooks

**Feature Store**

DASHBOARD

FEATURE LISTS

FEATURE SETS

Experiments

DATABASE

Data Import

SQL Editor

OLAP Console

Connect

PROCESS FLOWS

ETL Pipelines

MONITORING

Events









Logs

Activity

Help

## Feature Store Dashboard

### General


 Deployed Features 176	 Deployed Feature Set 2	 Training Views 3	 Training Sets 4
 Total Features 176	 Total Feature Set 2	 Associated Models 0	 Pending Feature Set 0

### Feature Usage



# End-to-End Governance

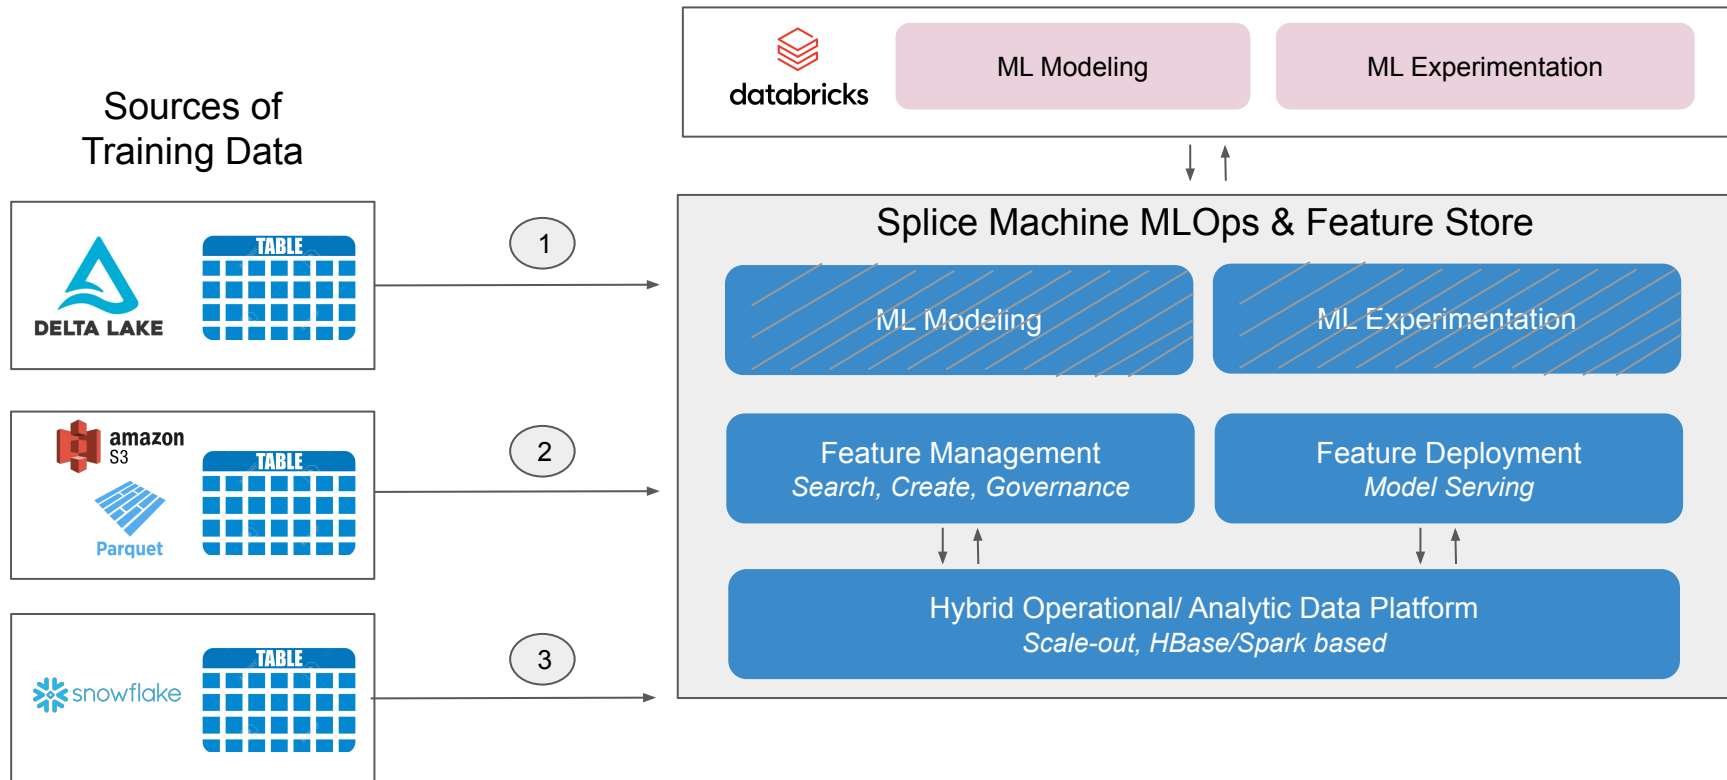
```
1 fs.display_model_feature_drift('deployed_models','retail_regression')
```



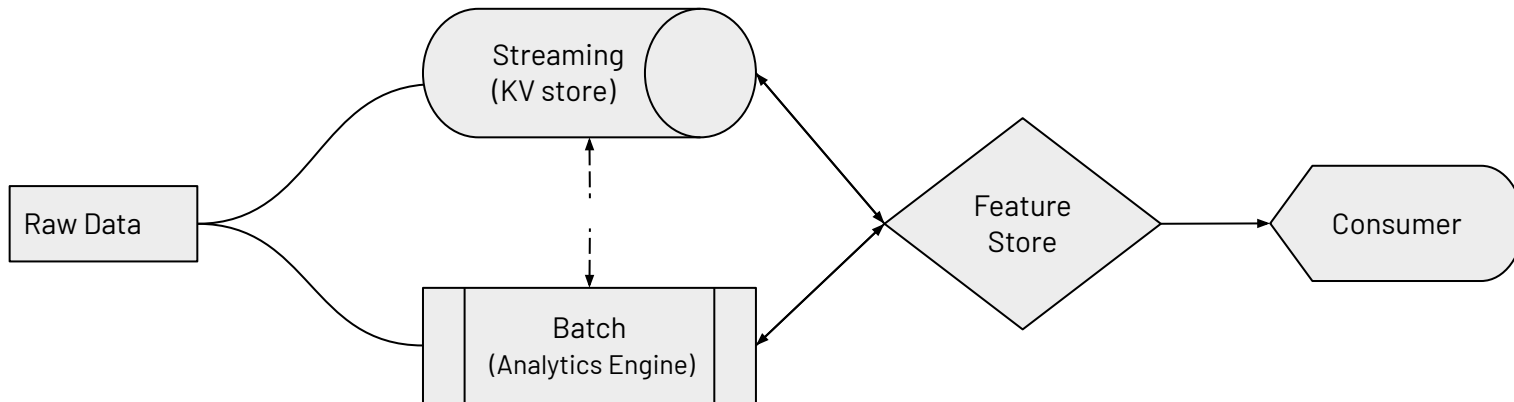
# How It Works



# How does this fit into the ML stack?



# Problem: Disconnected Compute Engines



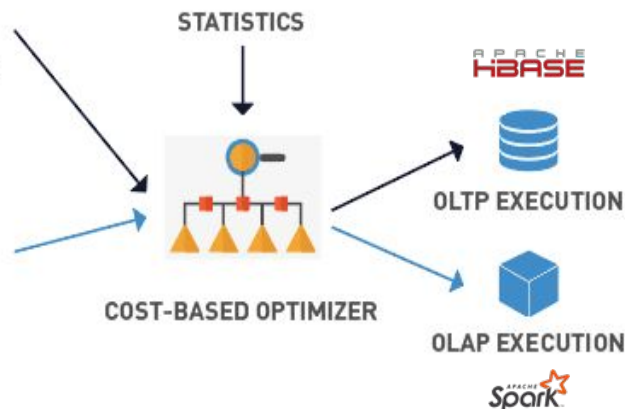
# Solution: Hybrid Operational/ Analytical RDBMS

- Scale-out
- Any Cloud/On-Prem
- Indexes and Triggers
- Full ACID Compliance

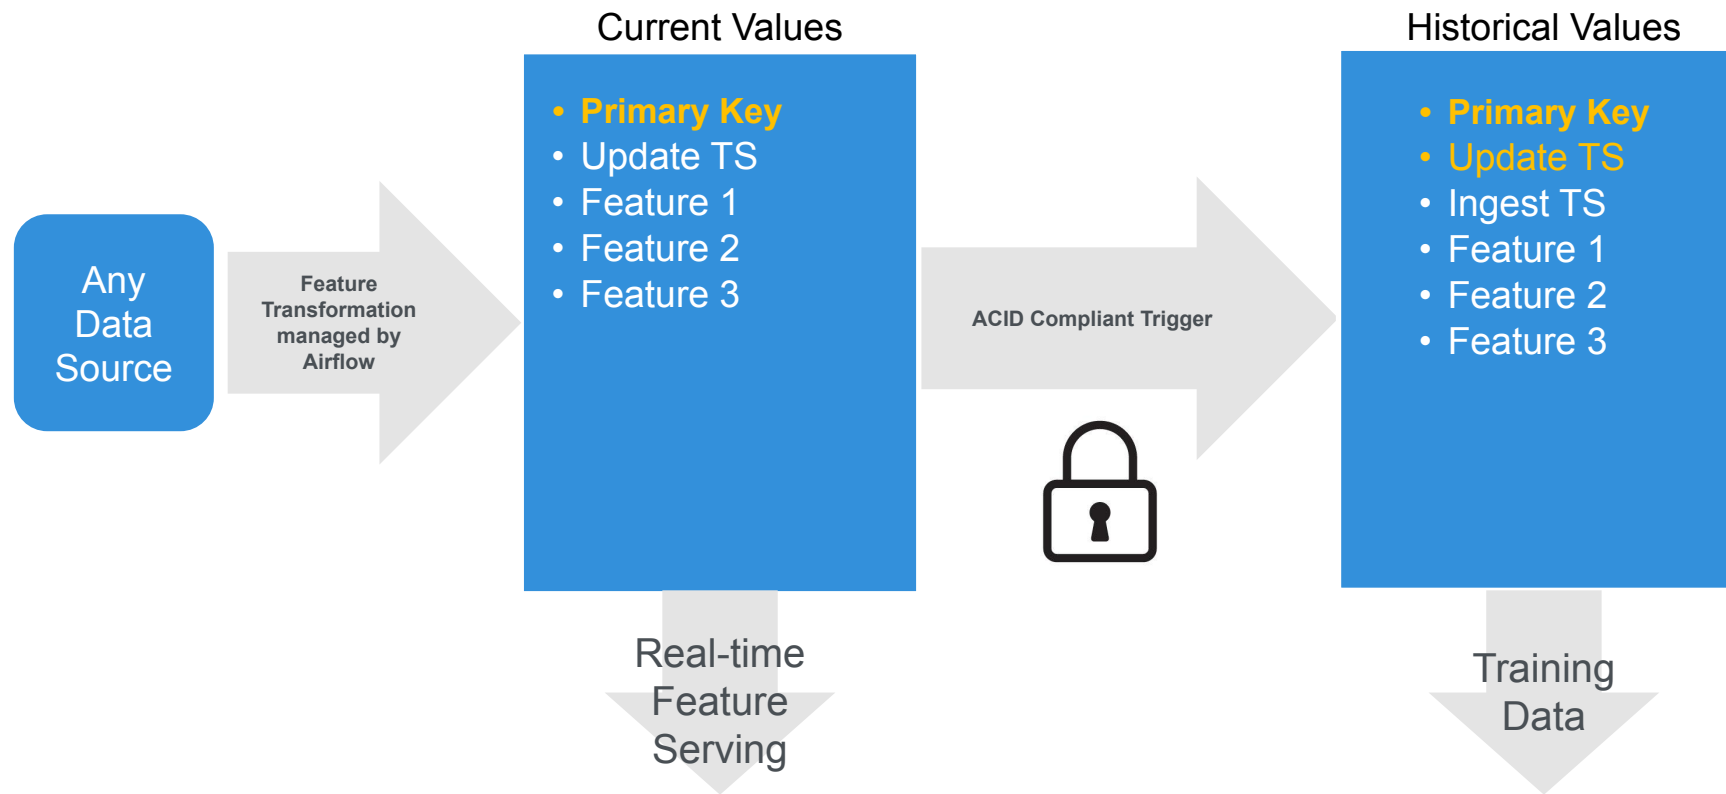
```
SELECT oli.order_id,  
       oli.order_lineitem_id,  
       oli.product_cost  
FROM orderlineitems oli  
WHERE oli.order_lineitem_id = 1004723876
```

```
SELECT sum(oli.product_cost),  
       EXTRACT(YEAR  
               FROM o.order_date) "Year",  
       EXTRACT(MONTH  
               FROM o.order_date) "Month",  
       cust.customer_name,  
       loc.location_name  
FROM orders o,  
     orderlineitems oli,  
     customer cust,  
     locations loc  
WHERE o.order_id = oli.order_id  
      AND o.customer_id = cust.customer_id  
      AND o.location_id = loc.location_id  
      AND o.order_date > '2018-01-01'  
GROUP BY EXTRACT(YEAR  
                 FROM o.order_date),  
         EXTRACT(MONTH  
                 FROM o.order_date),  
         cust.customer_name,  
         loc.location_name
```

SQL QUERIES



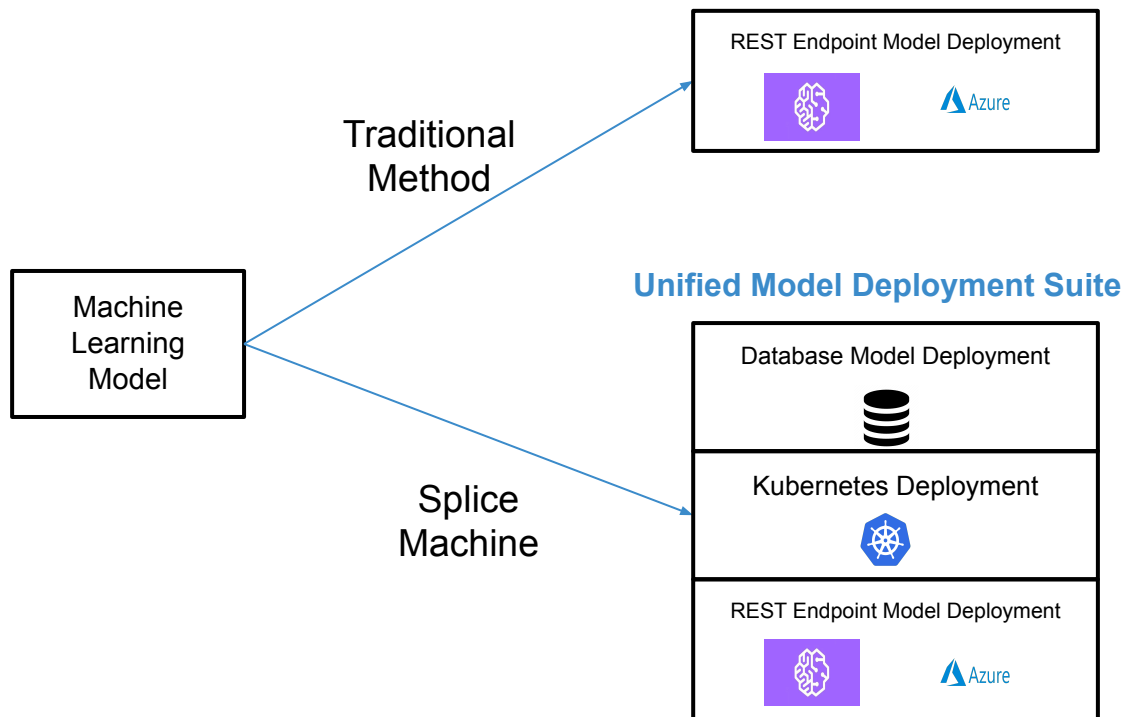
# Feature Set- Two Tables in a Single Database



# MLOps



- 1-line deployment
- Deploy models to database
- Features and prediction are memorialized



# Database Deployment

- New records automatically trigger predictions
- Single click - low code deployment
- Easy model governance and traceability via SQL in an “evaluation store”
- Sub-millisecond predictions and scalable
- No extra endpoint programming

User	TS	RunID	F1	F2	F3	Pred
Jack	2020-09-01	4400	5.2	-12	15.4	Fraud
John	2020-09-01	4401	9.4	4	2.3	No Fraud
Sarah	2020-09-25	4402	2.3	7	1.1	Fraud
Steve	2020-10-12	4403	1.3	0	3.4	...



Prediction made and populated at sub-millisecond speed

# Benefits



Data Preparation

Feature Engineering

Experimentation

Deployment



Data  
Preparation

Feature  
Engineering

Experimentation

Deployment

Automated  
Aggregation  
Pipelines

`ffs.get_feature_set_from_source`

Feature Reuse  
`ffs.serach_feature_store`


Automated  
Training Set  
`fs.get_training_set_from_view`

Database  
Deployment  
`ml.deploy_model`

100X Faster


# Coming in 2021

- Comprehensive how-to text
- Hands-on exercises
- Best practices





# Summary



## Improve Productivity

Feature engineering is streamlined and automated



## Predictive Accuracy

Training and serving pipelines are consistent and repeatable



## Model Governance

Full transparency

# Next Steps

Get Demo: <https://www.splicemachine.com>

Free Trial: <https://cloud.splicemachine.io/login>

We are hiring!

# Q&A