



How Facebook Brought QUIC To Billions

An Engineering Journey



Matt Joras

Software Engineer, Traffic Protocols
Co-chair IETF QUIC WG
mjoras@fb.com



Yang Chi

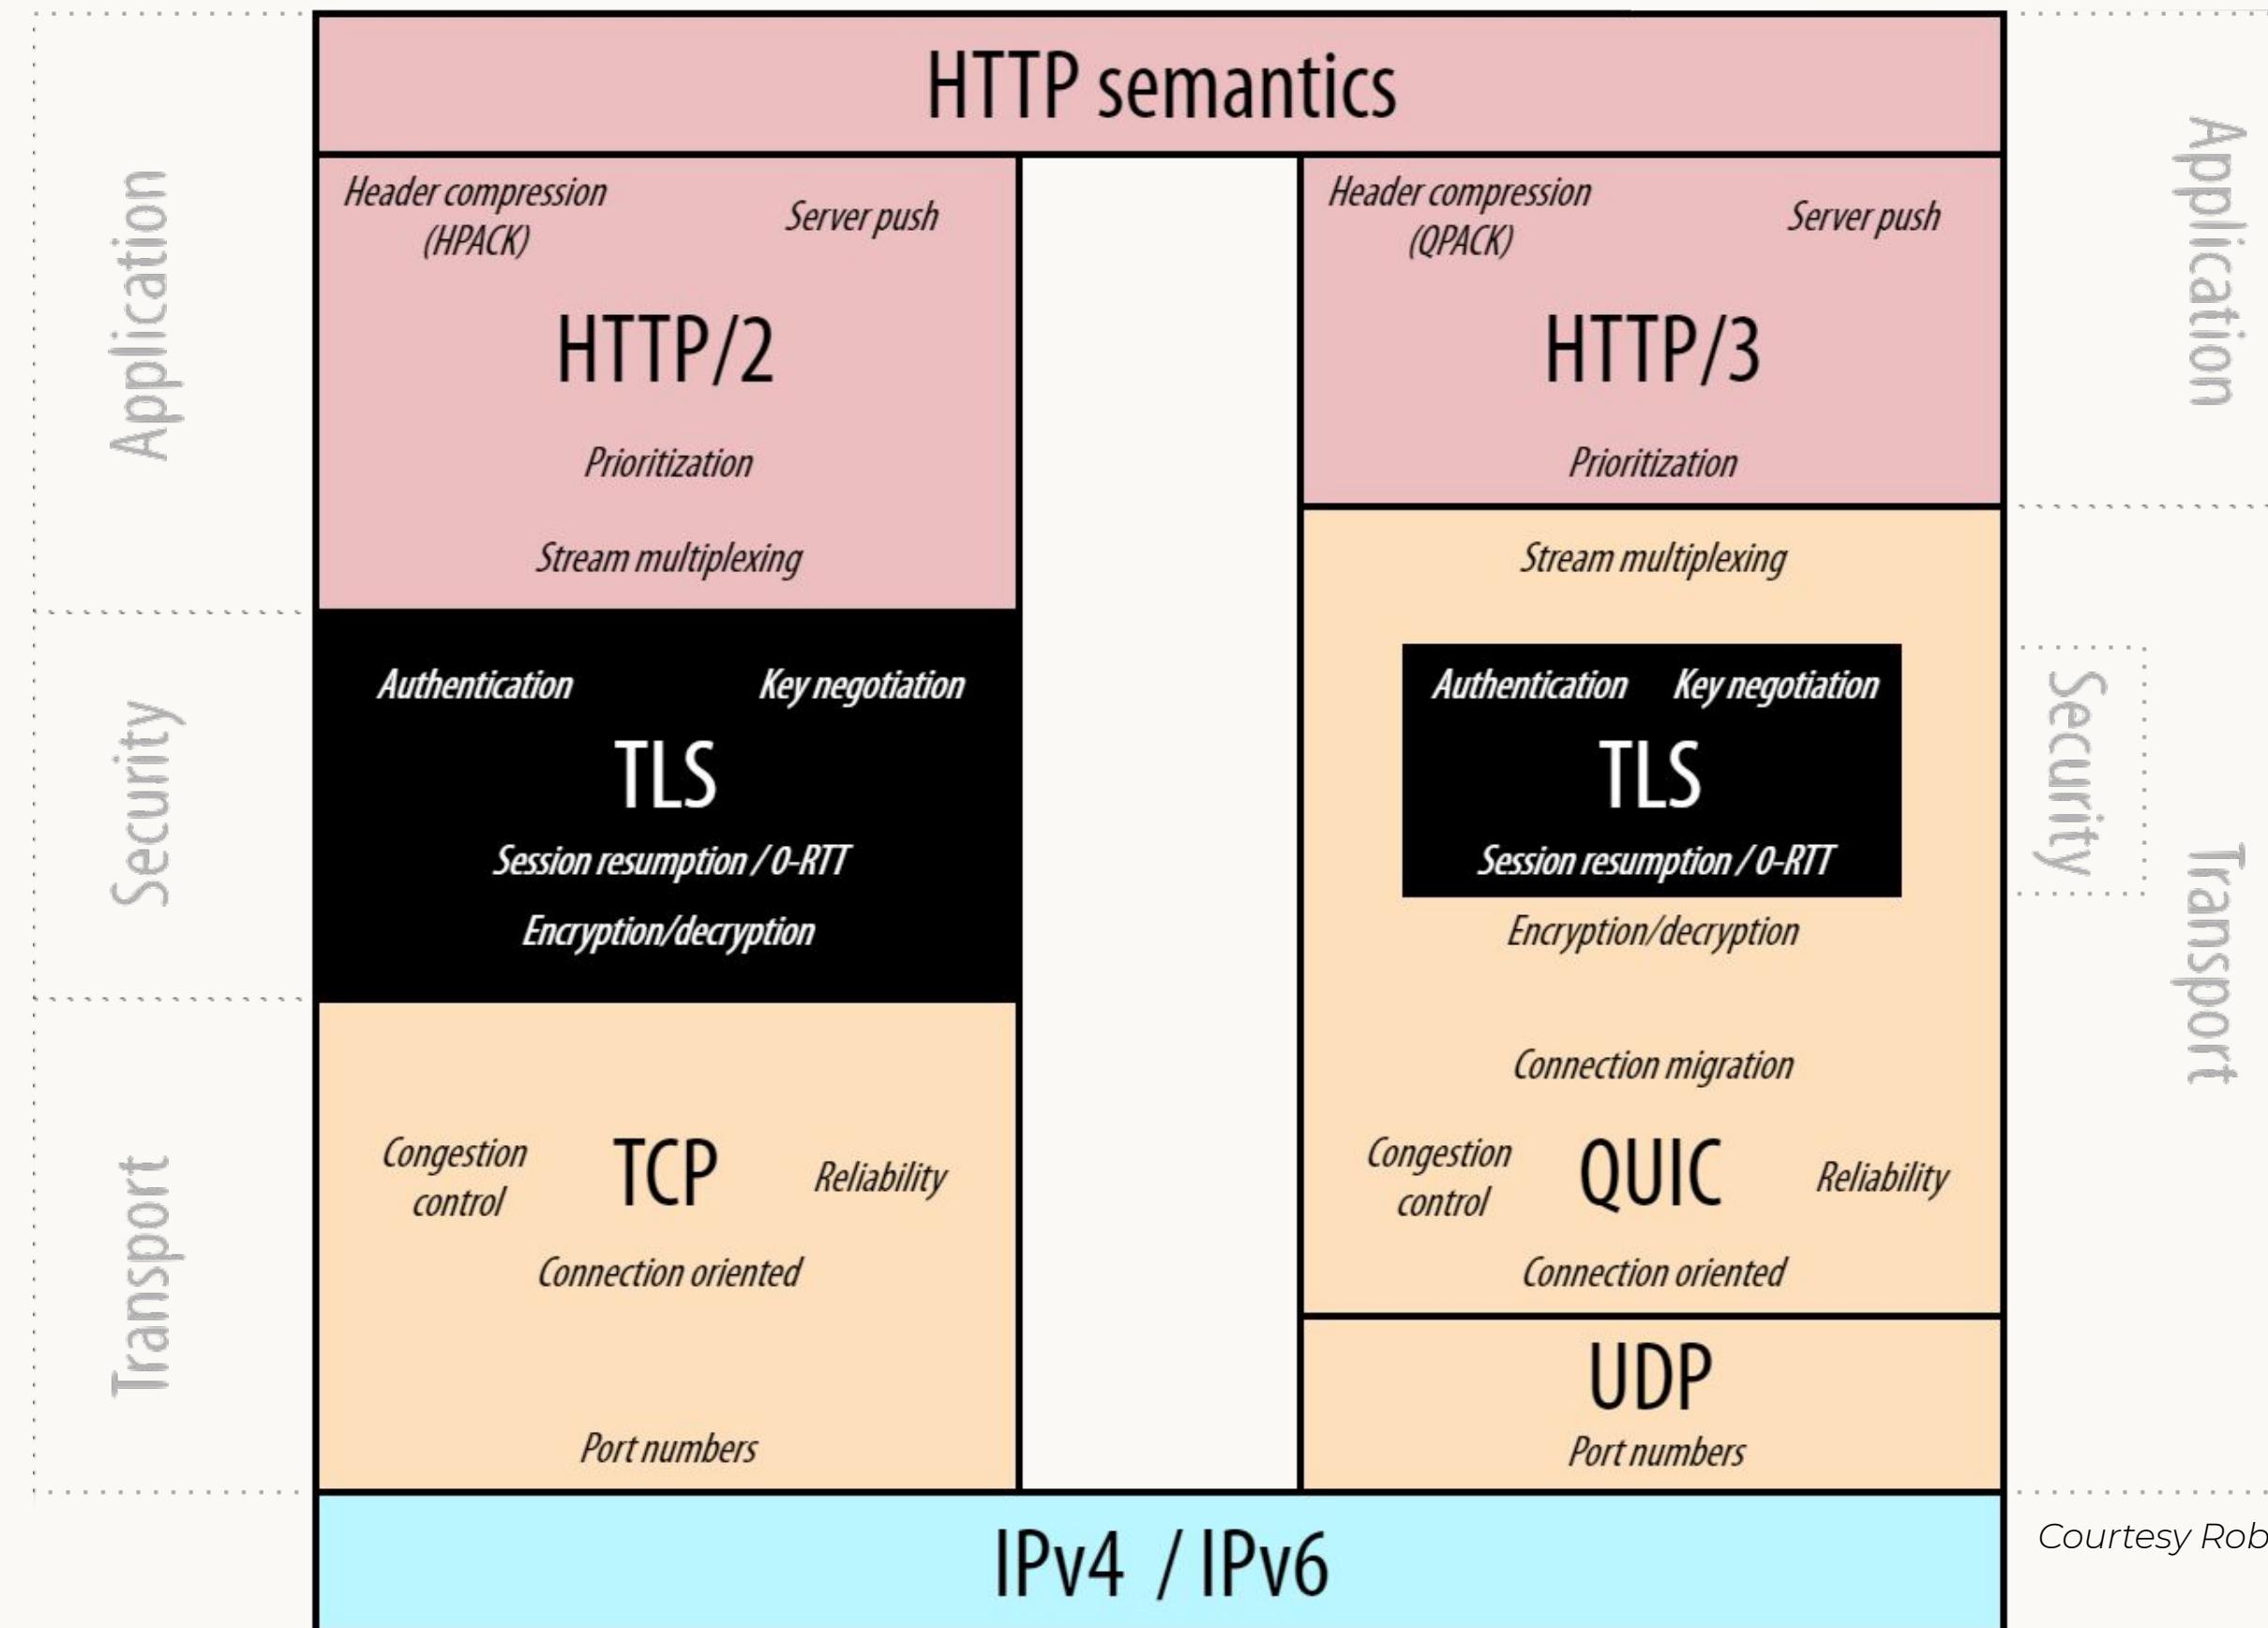
Software Engineer, Traffic Protocols
yangchi@fb.com

What is QUIC? The next Internet transport protocol.

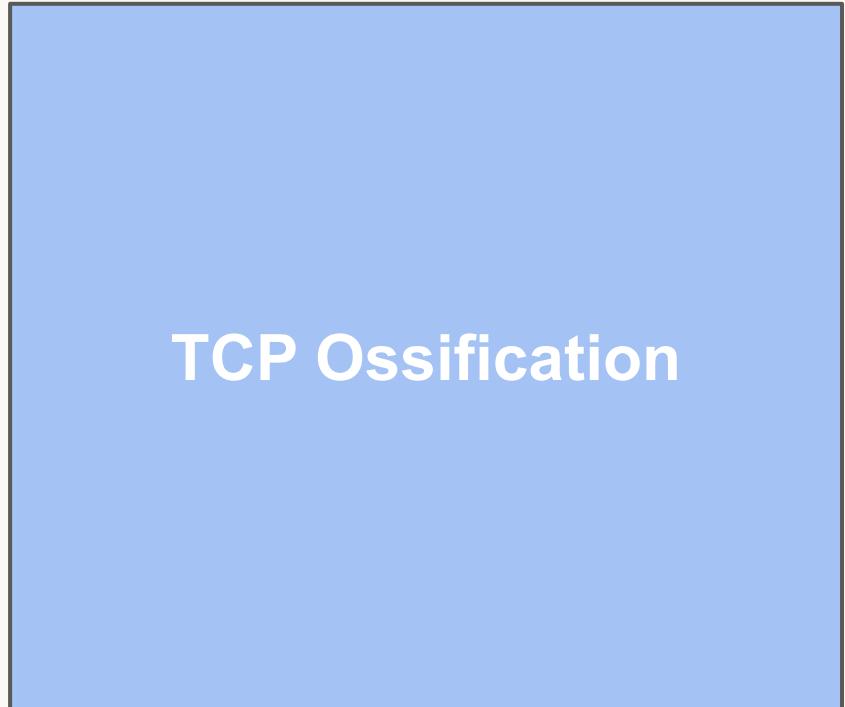


QUIC

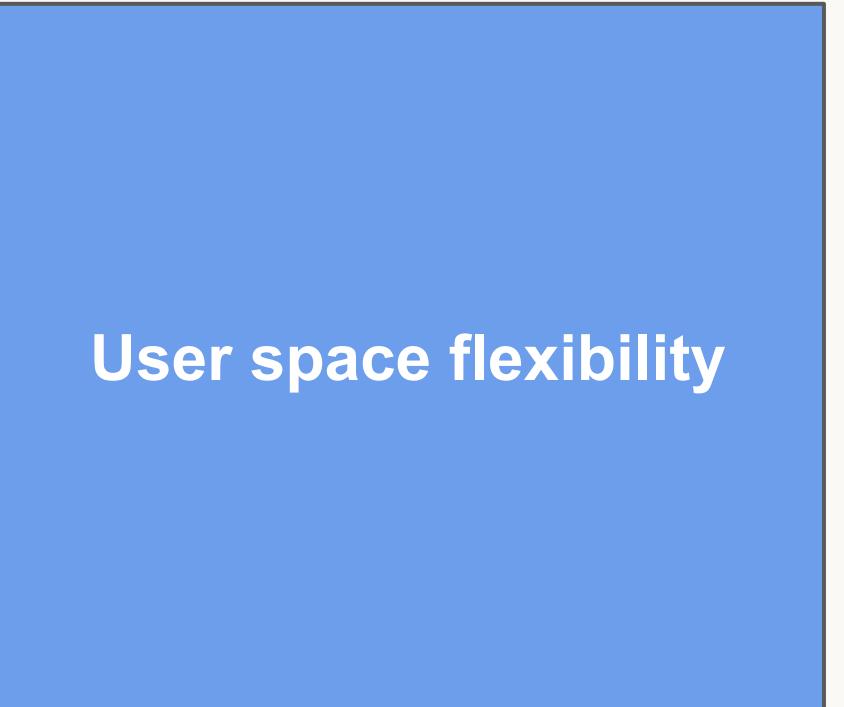
What is QUIC?



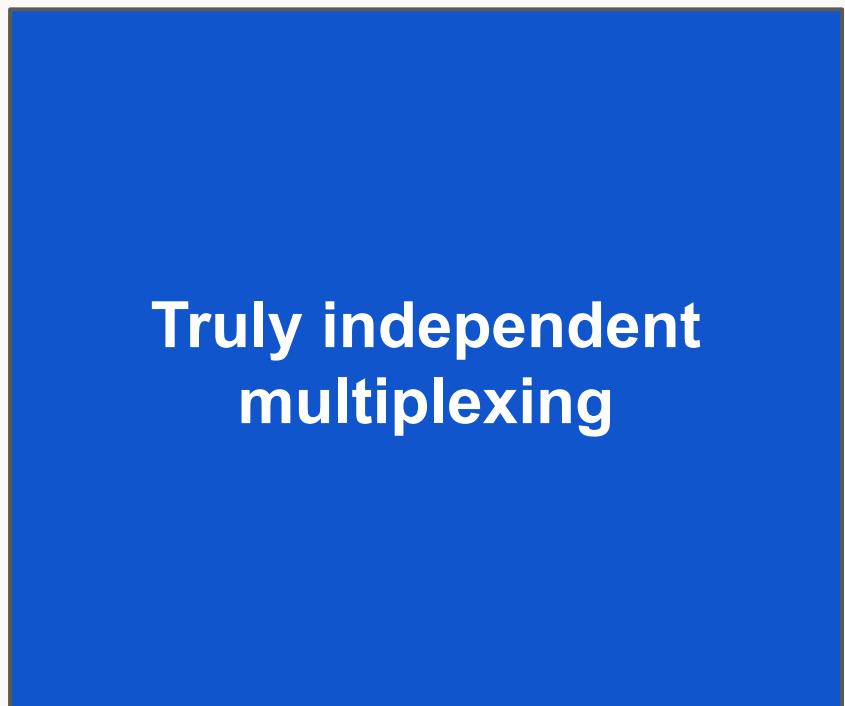
Why QUIC?



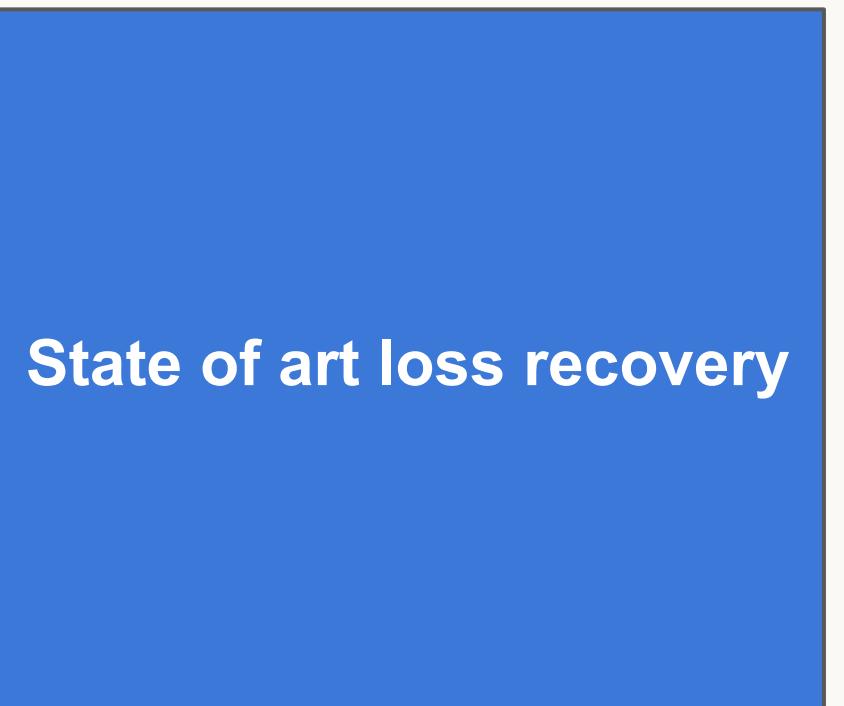
TCP Ossification



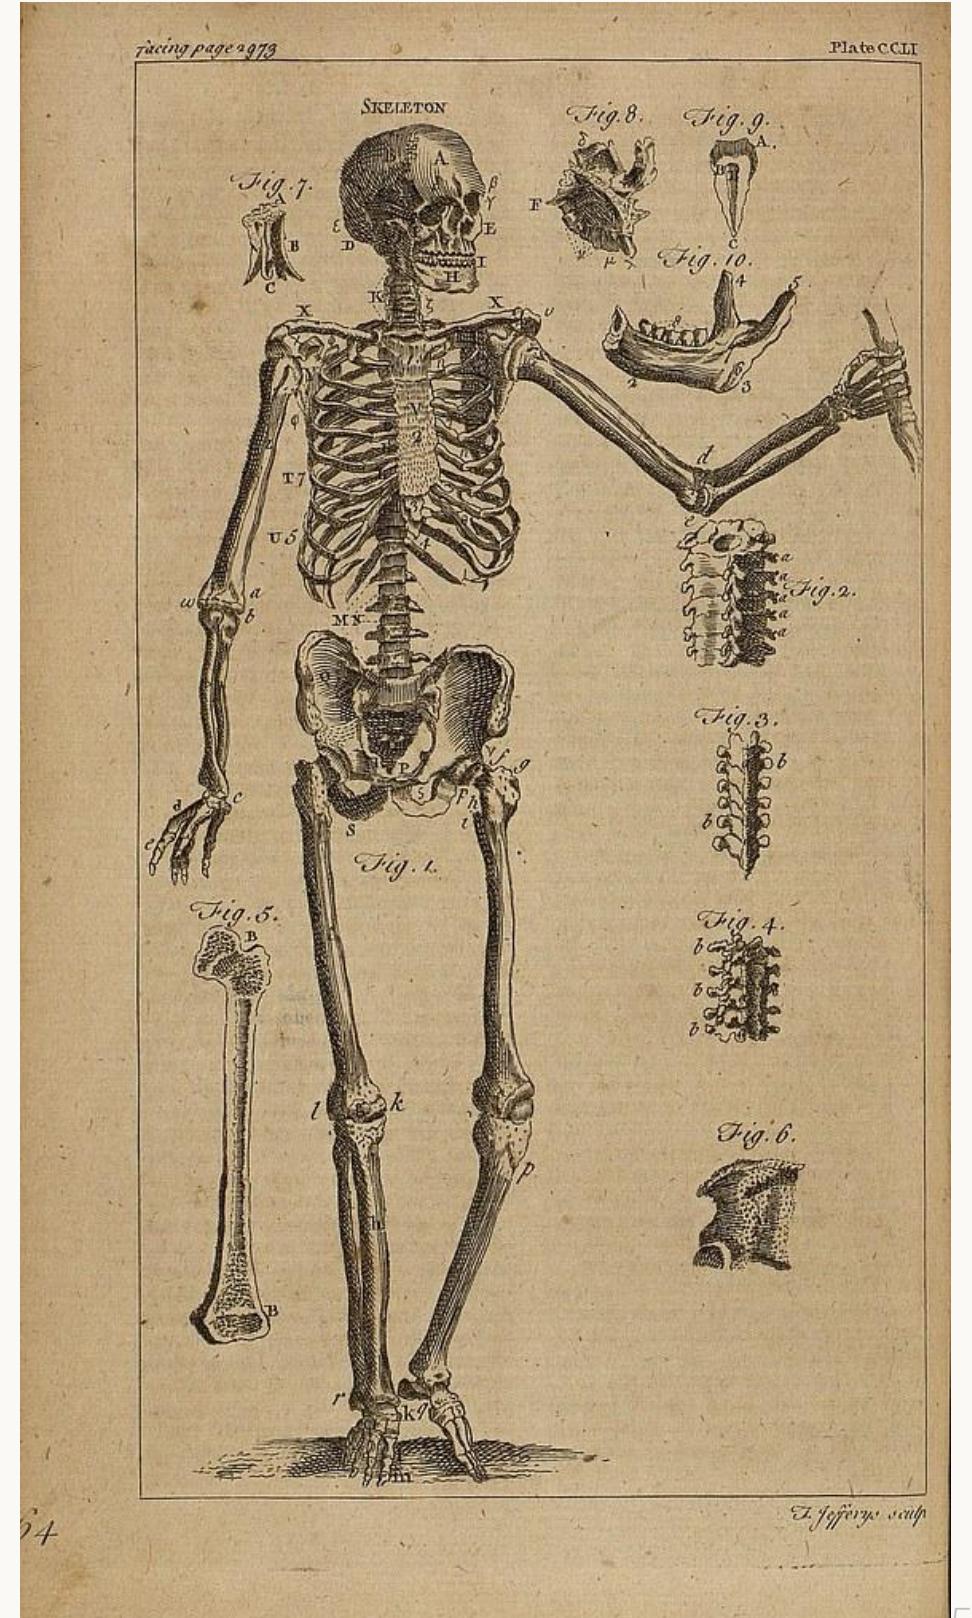
User space flexibility



Truly independent
multiplexing



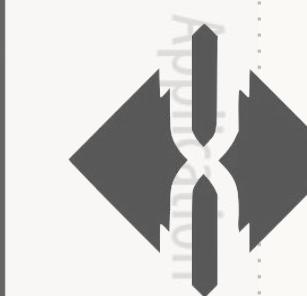
State of art loss recovery



QUIC@FB

Application Security Transport

HTTP semantics		
Header compression (HPACK)	Server push	Header compression (QPACK)
HTTP/2	Prioritization	HTTP/3
Stream multiplexing		Stream multiplexing
Authentication	Key negotiation	Authentication
TLS		TLS
Session resumption / 0-RTT		Session resumption / 0-RTT
Encryption/decryption		Encryption/decryption
Congestion control	TCP	Reliability
Connection oriented		Connection migration
Port numbers		Port numbers
IPv4 / IPv6		



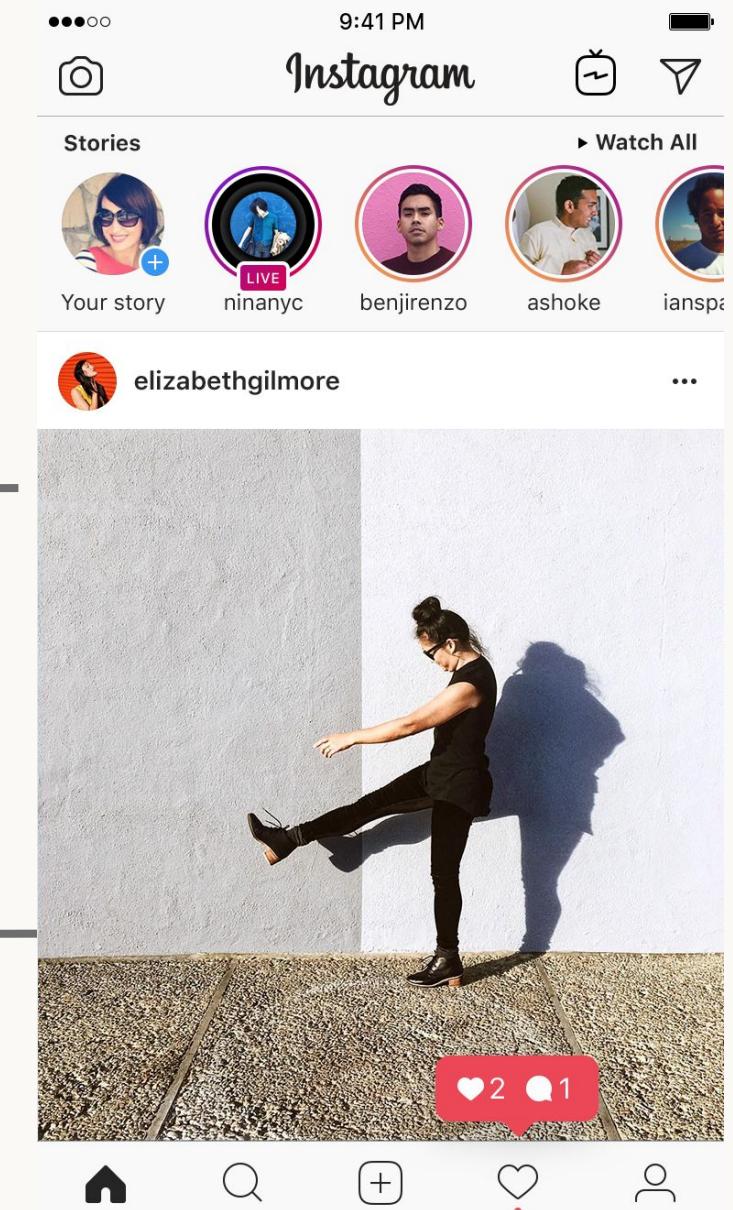
Proxygen



Fizz

MFVfst

Katran





Sounds great. So you just turned on
QUIC and everything got better?





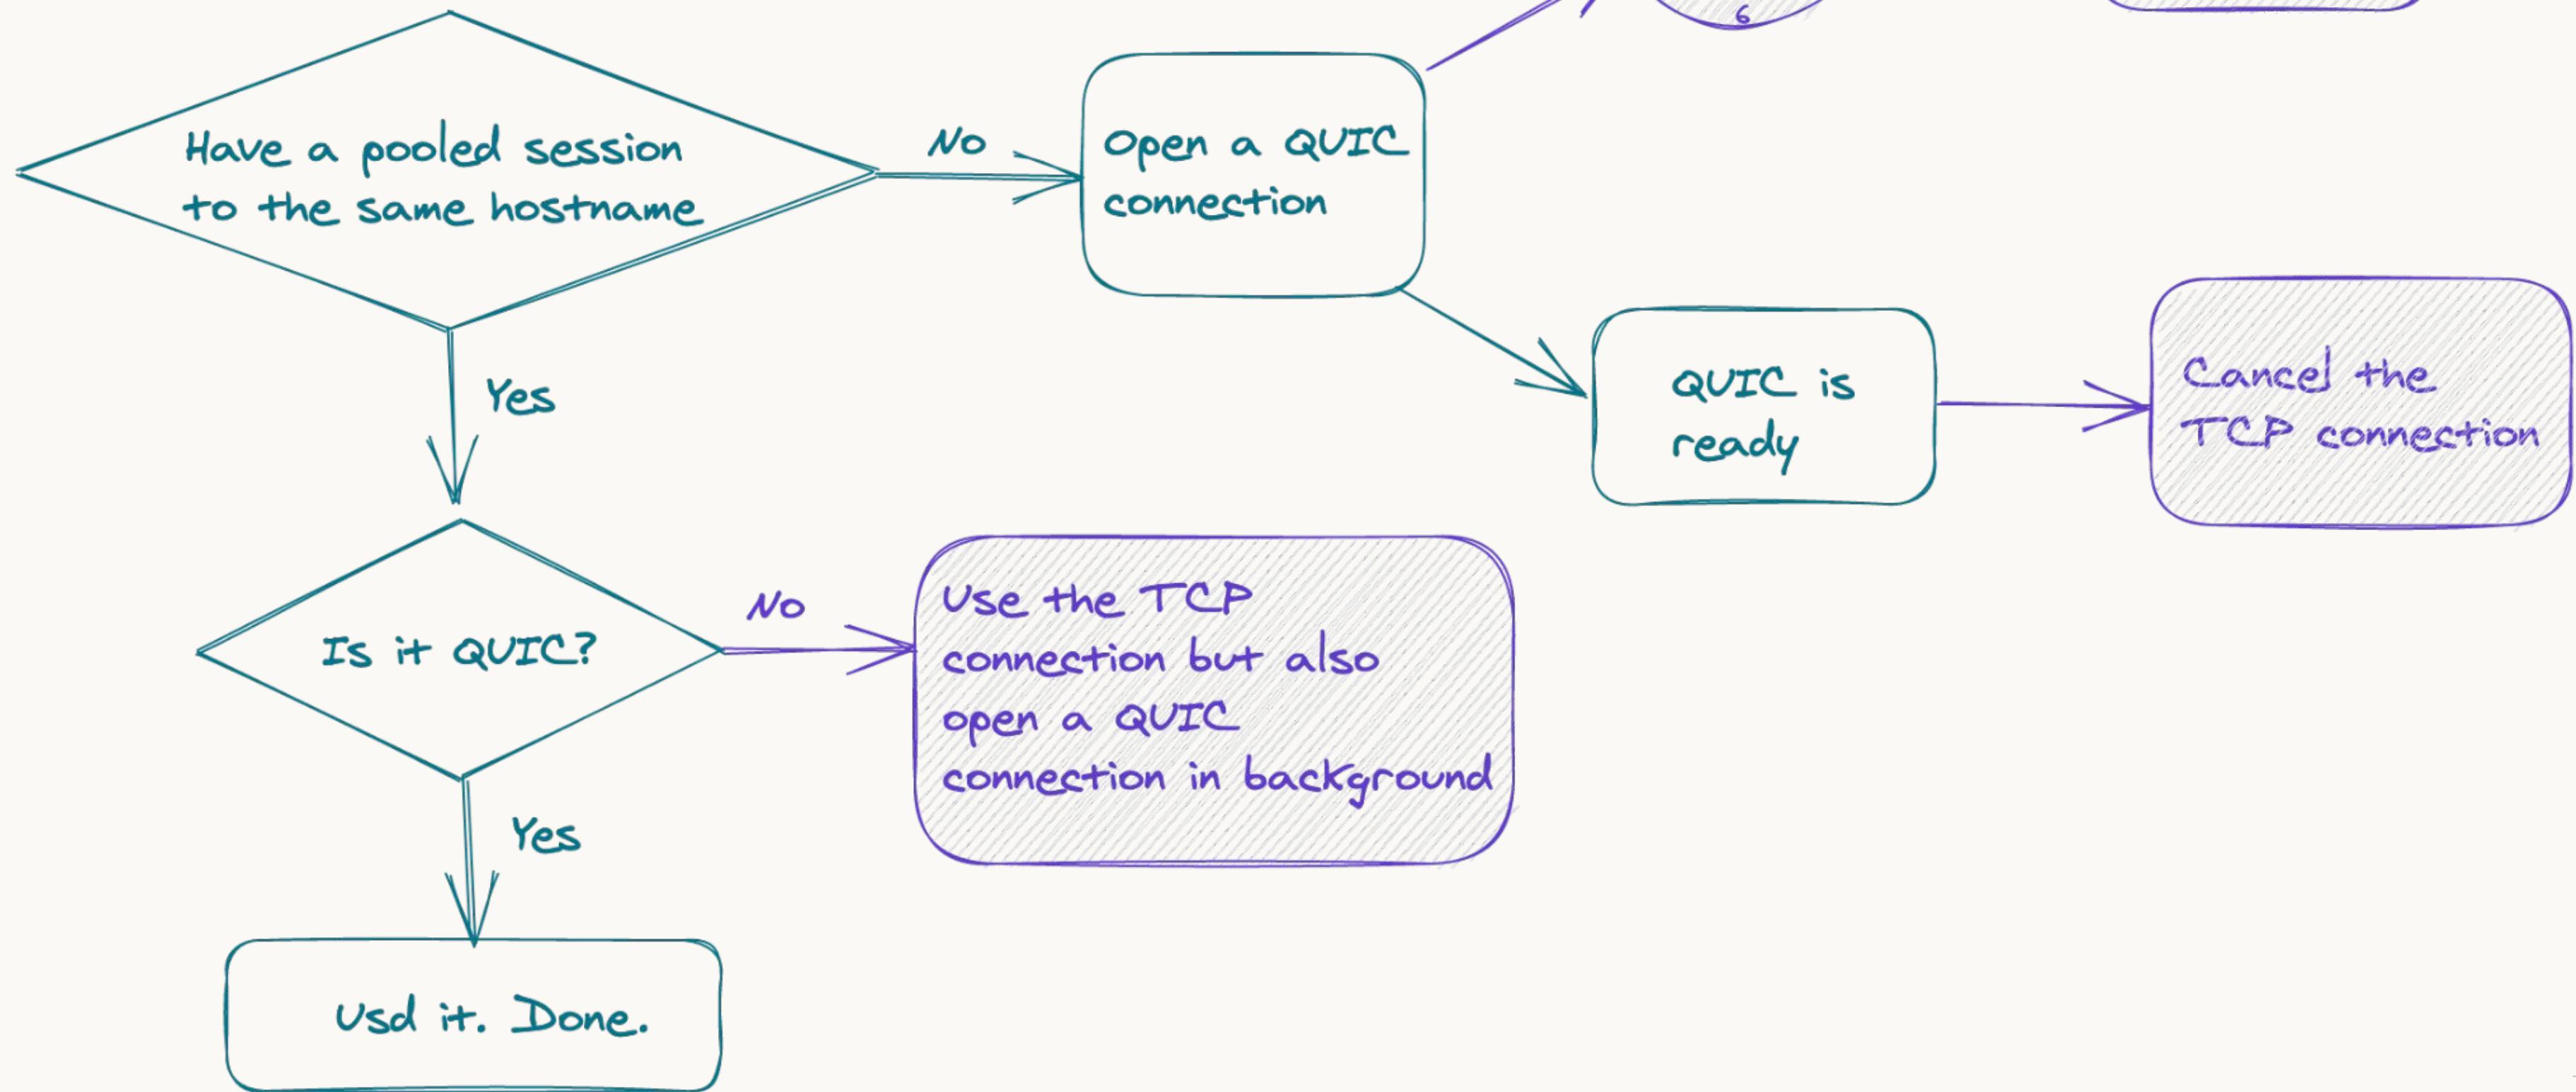
Mobile apps required significant investigation to iron out issues with QUIC.

In several cases, app behavior and some settings were implicitly tuned for TCP.

FB & IG Baselines

		
Transport	TCP + TLS 1.3	TCP + TLS 1.3
GraphQL / API requests	Single HTTP/2	Single HTTP/2
Image requests	Single HTTP/2	Up to 6 HTTP/1.1 connections
Video requests	Single HTTP/2	Up to 6 HTTP/1.1 connections
Connection sharding	GraphQL, Images and videos cannot share connections	Images and videos can use the same connections. API is over separate connection.
Server side stack	Vanilla Linux TCP + BBRv1	Vanilla Linux TCP + BBRv1

QUIC with TCP backup



Congestion Control and Flow Control

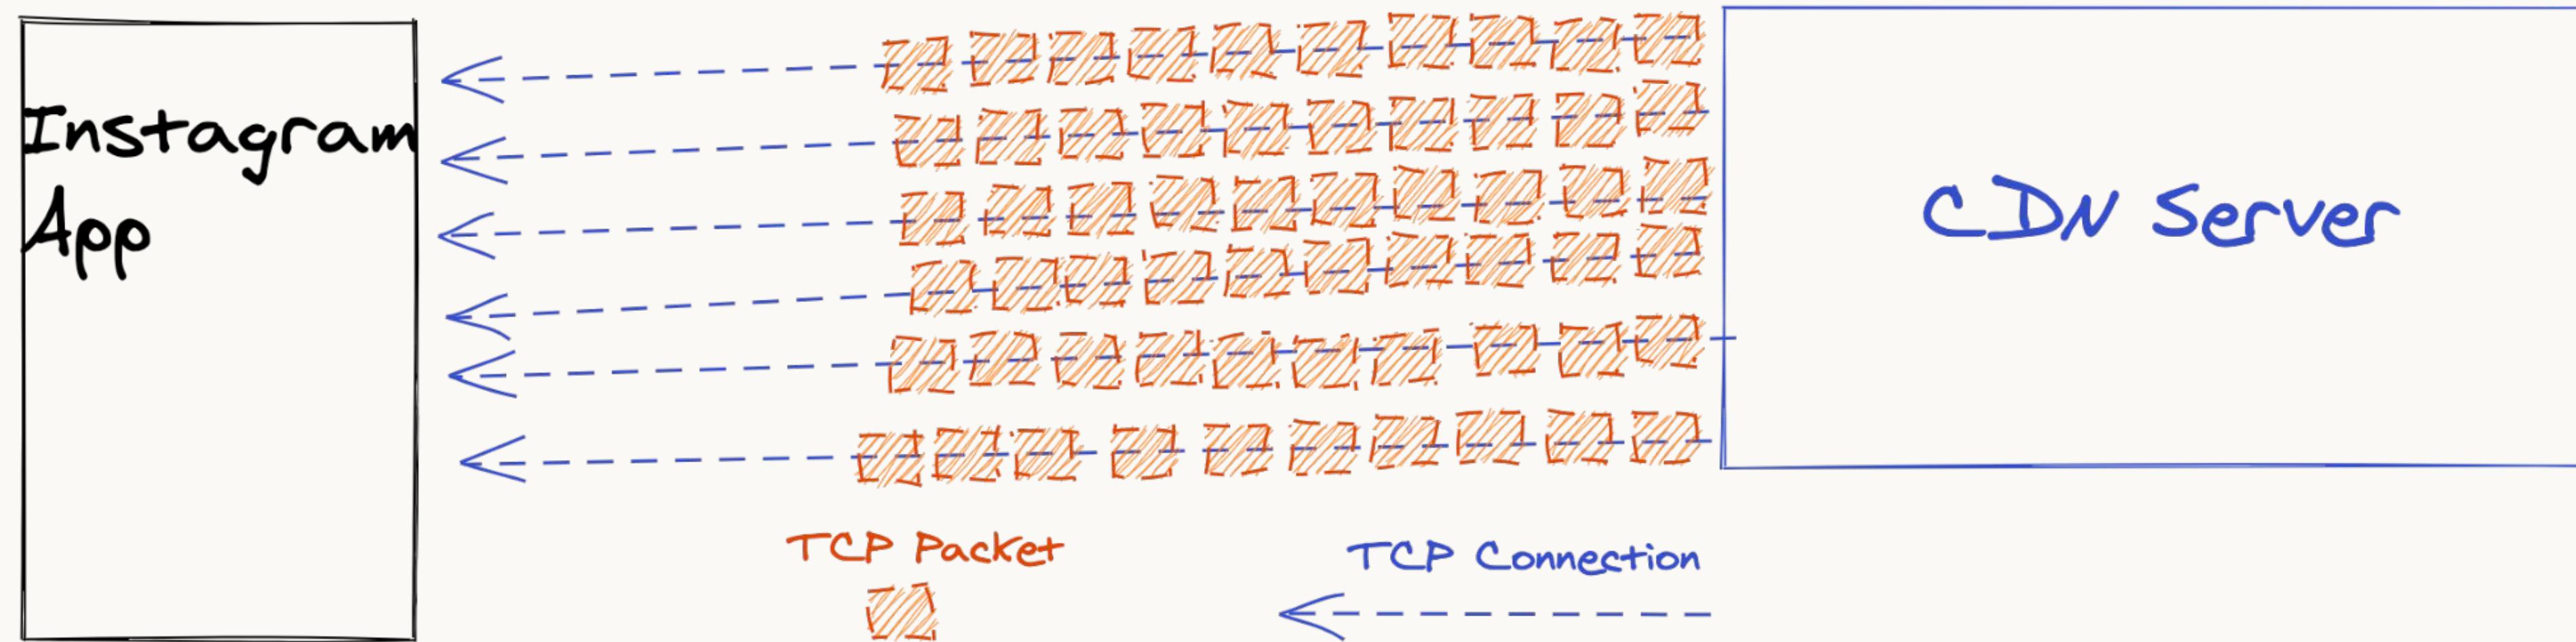
Congestion
Control

Algorithm used to minimize loss and maximize throughput.

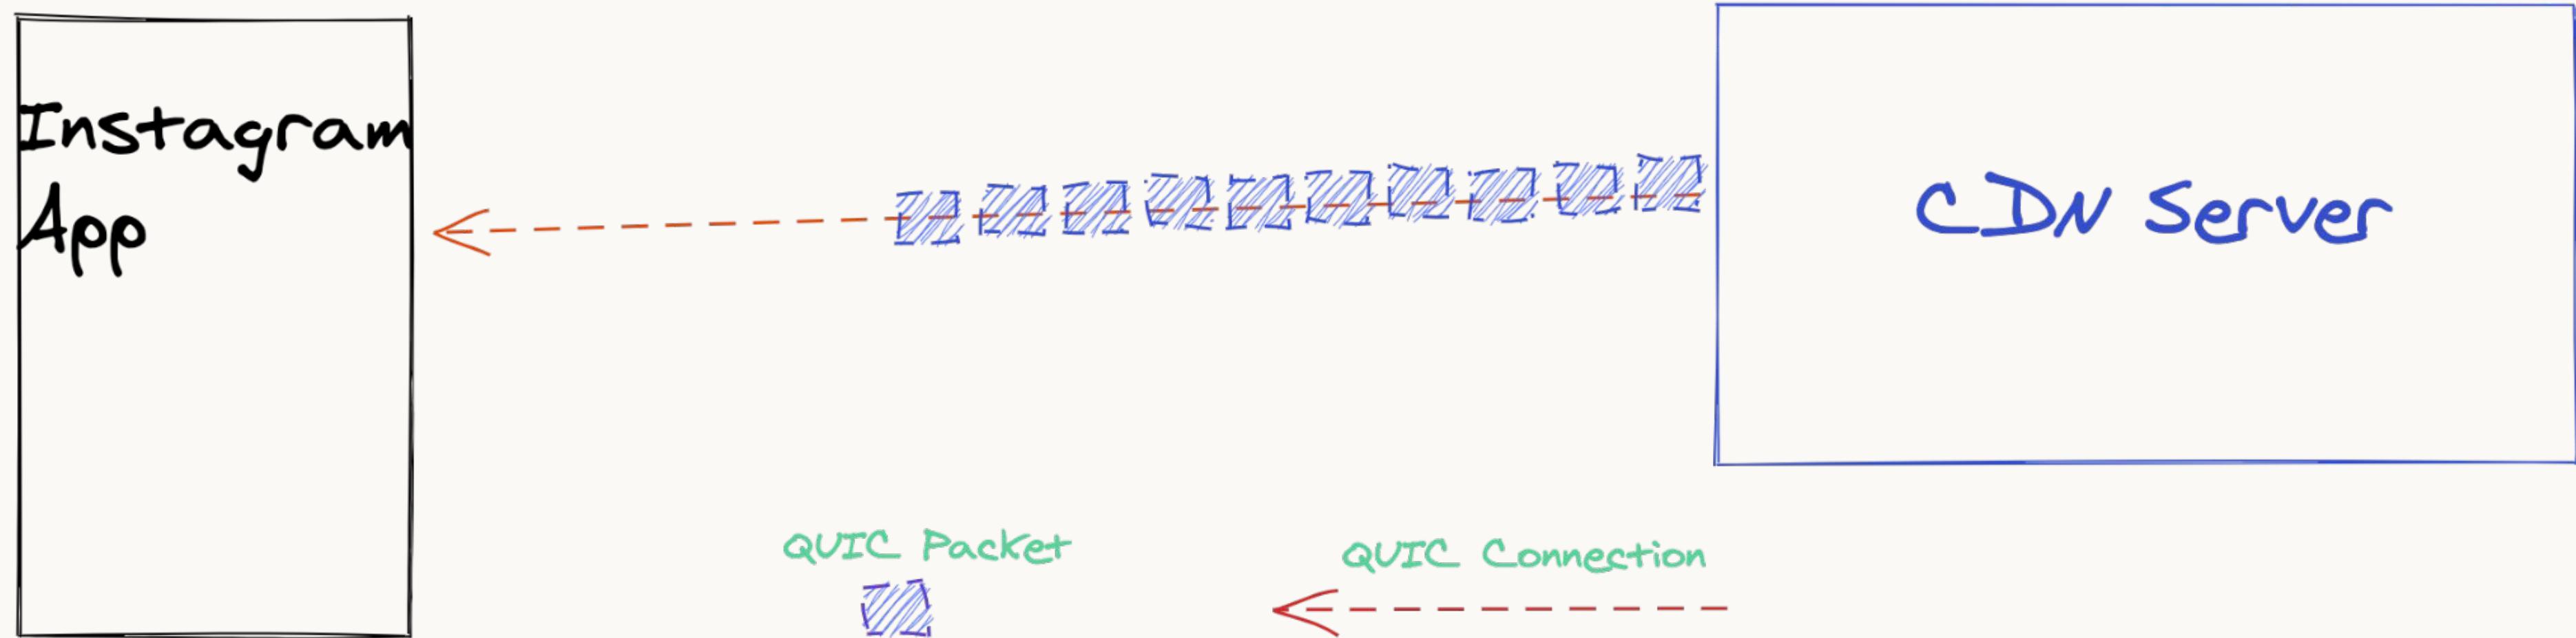
Flow Control

Limits the amount of data the applications buffers.

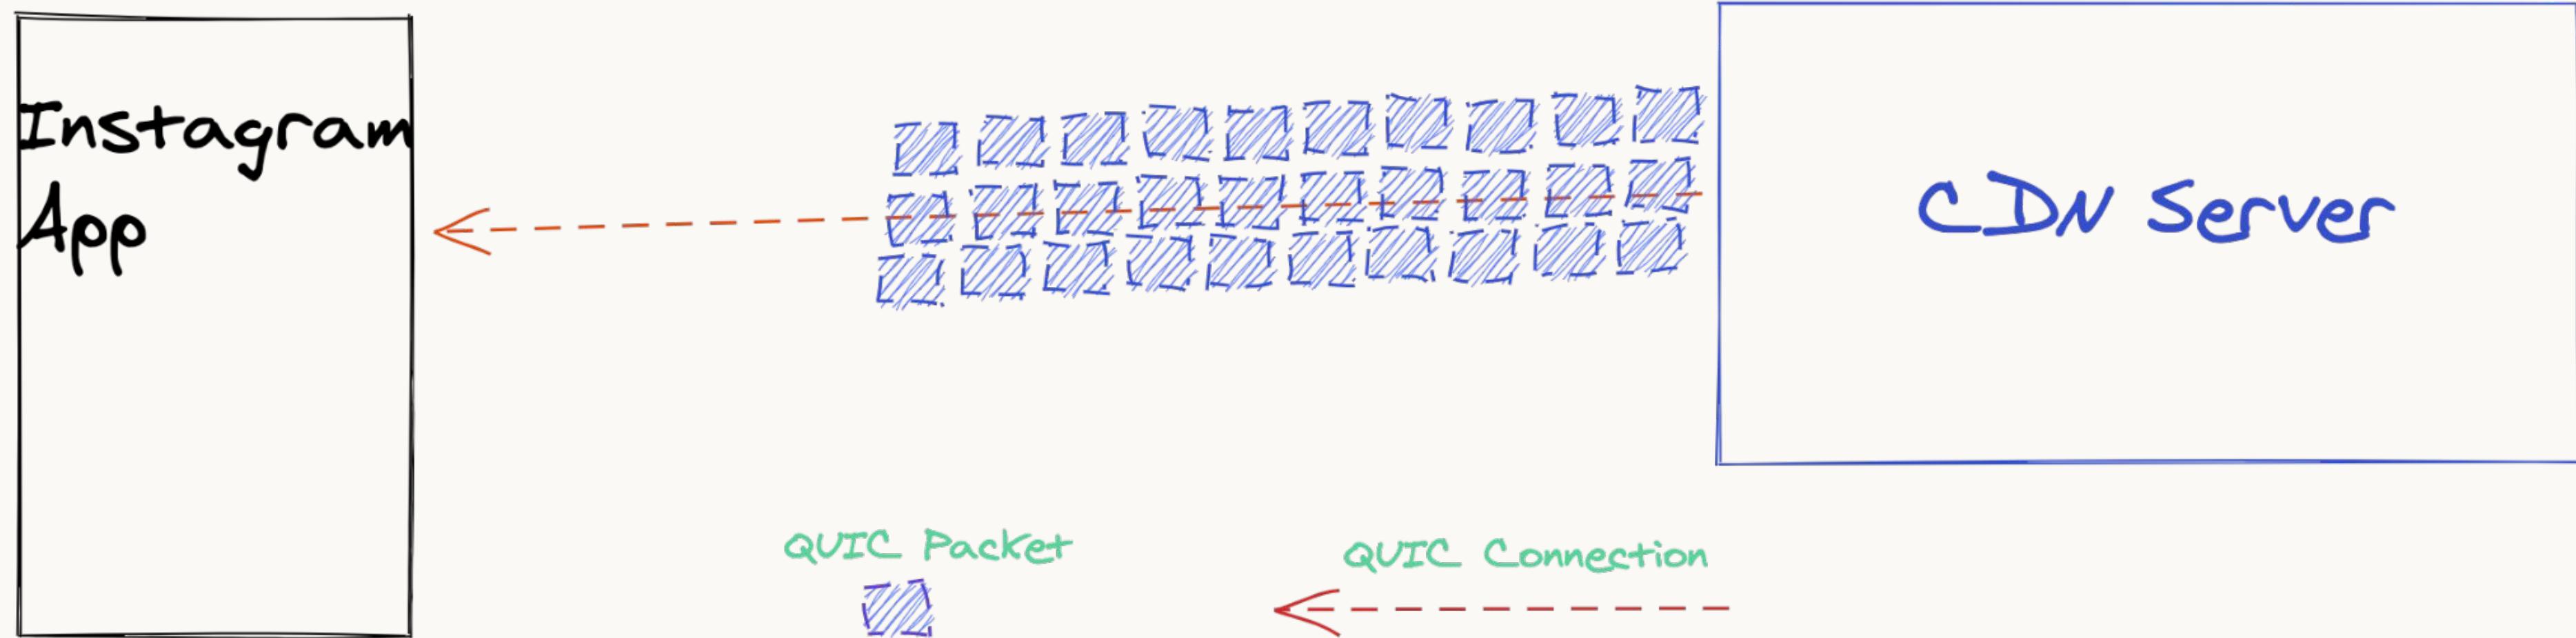
Initial flight: 6 TCP conns vs 1 QUIC conn



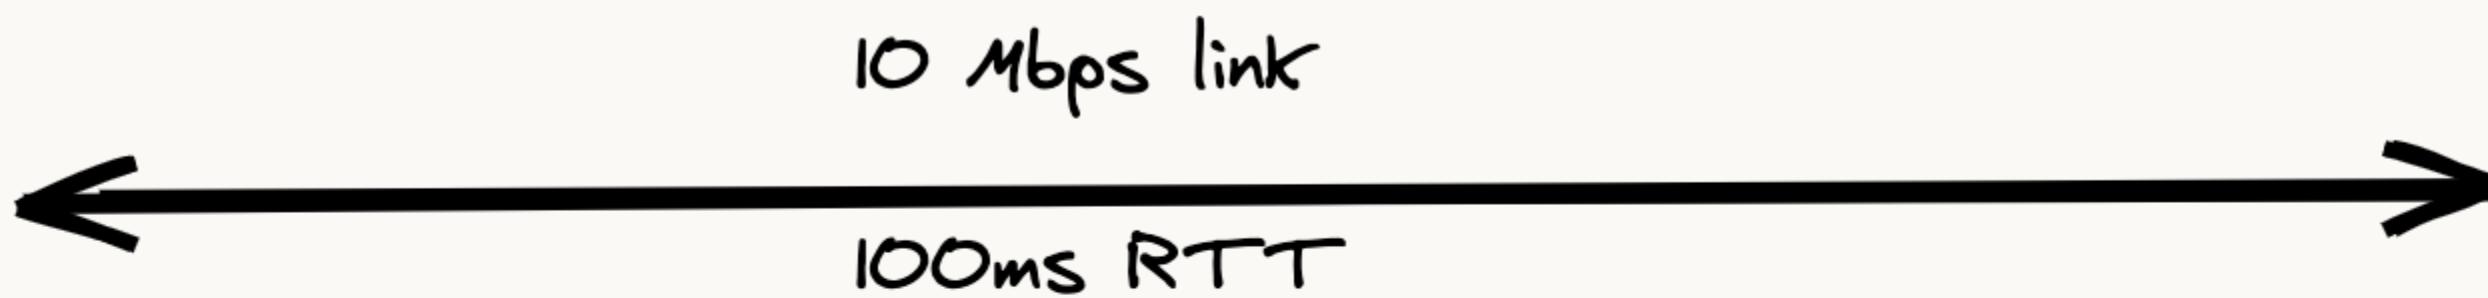
Initial flight: 6 TCP conns vs 1 QUIC conn



Initial flight: 6 TCP conns vs 1 QUIC conn

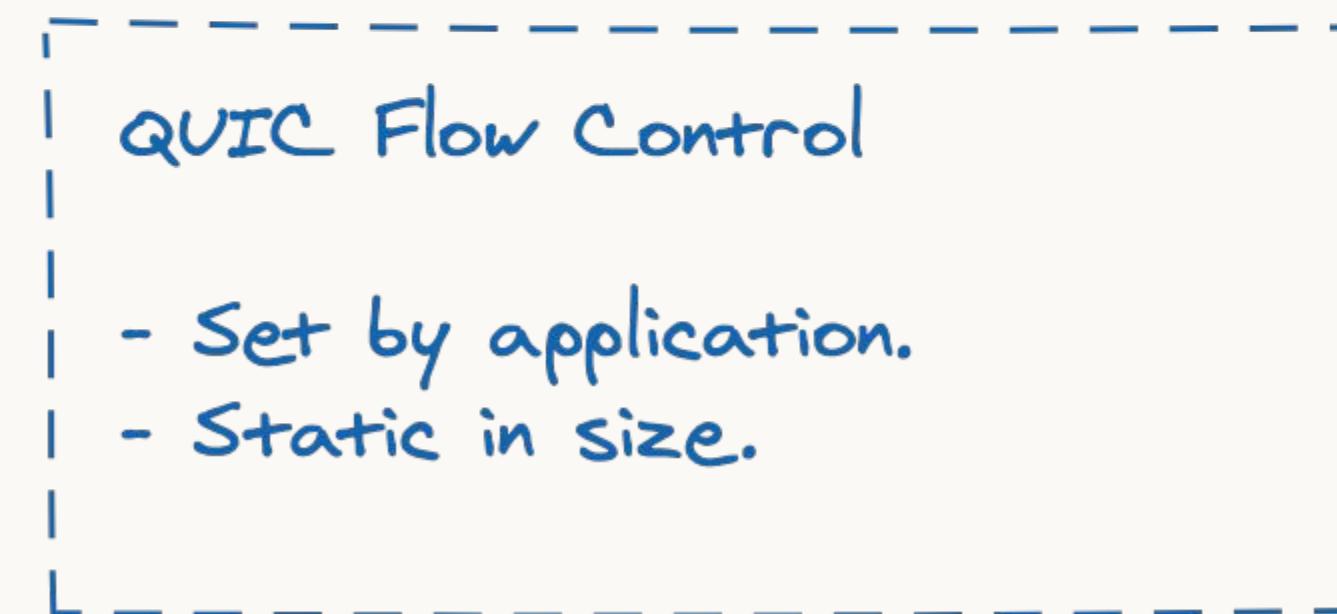


Congestion Control and Congestion Window

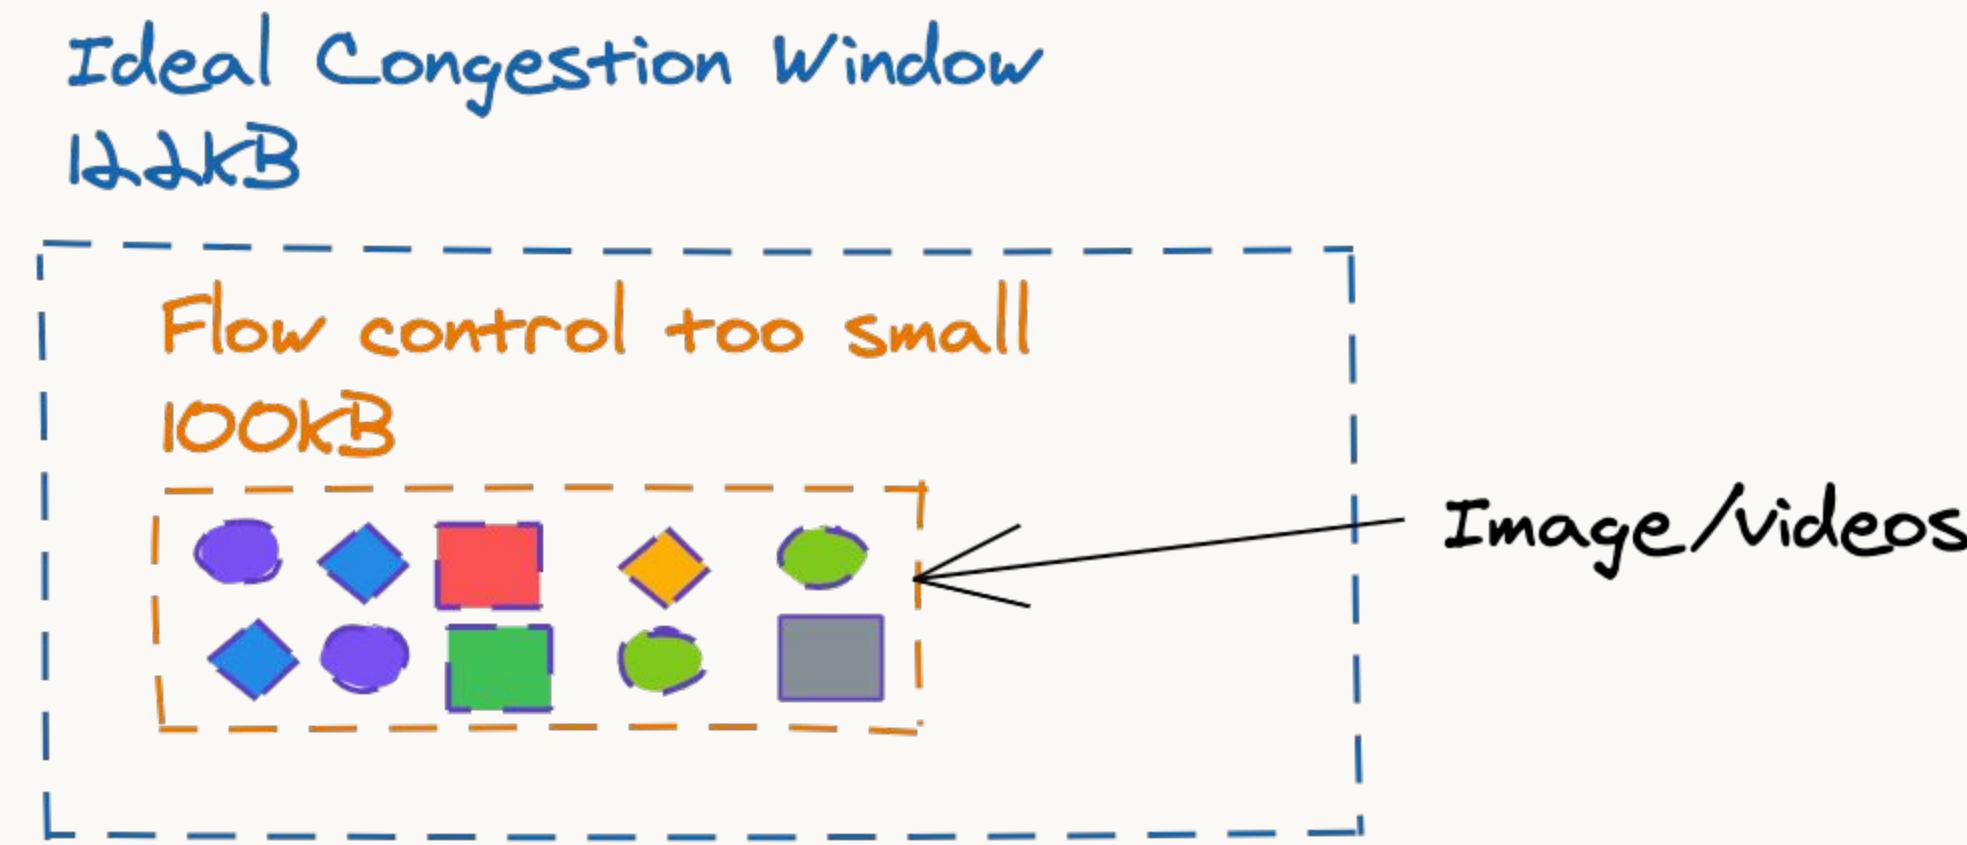


bandwidth * delay = data needed to achieve full throughput
 $10\text{Mbps} * 100\text{ms} = 122\text{kB}$ = ideal "congestion window"

TCP + HTTP/2 Has Two Levels of Flow Control



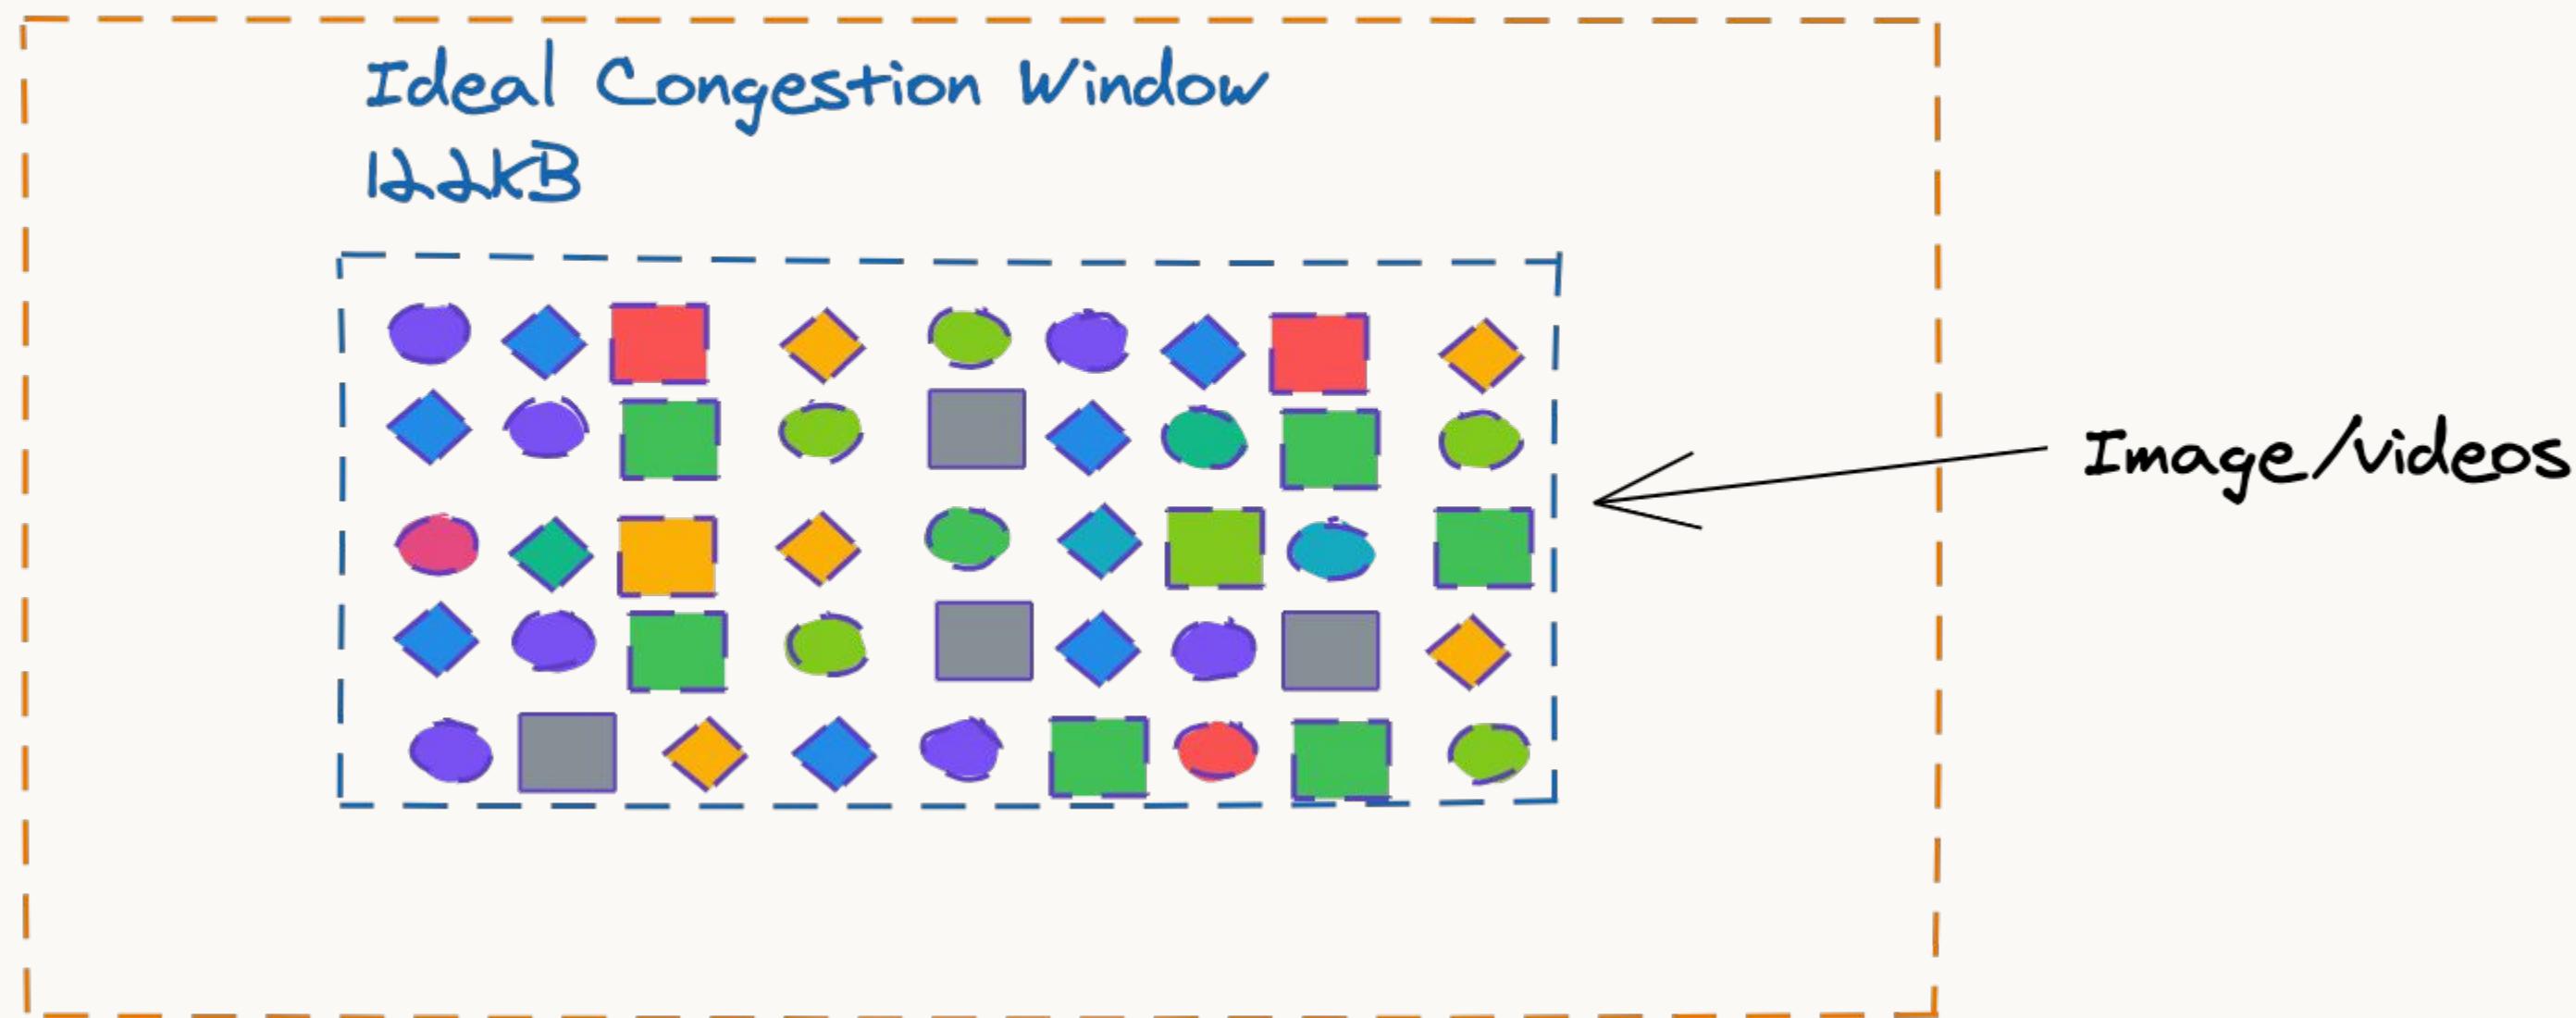
Flow Control Constrains Bandwidth



Increasing Flow Control Solves The Problem

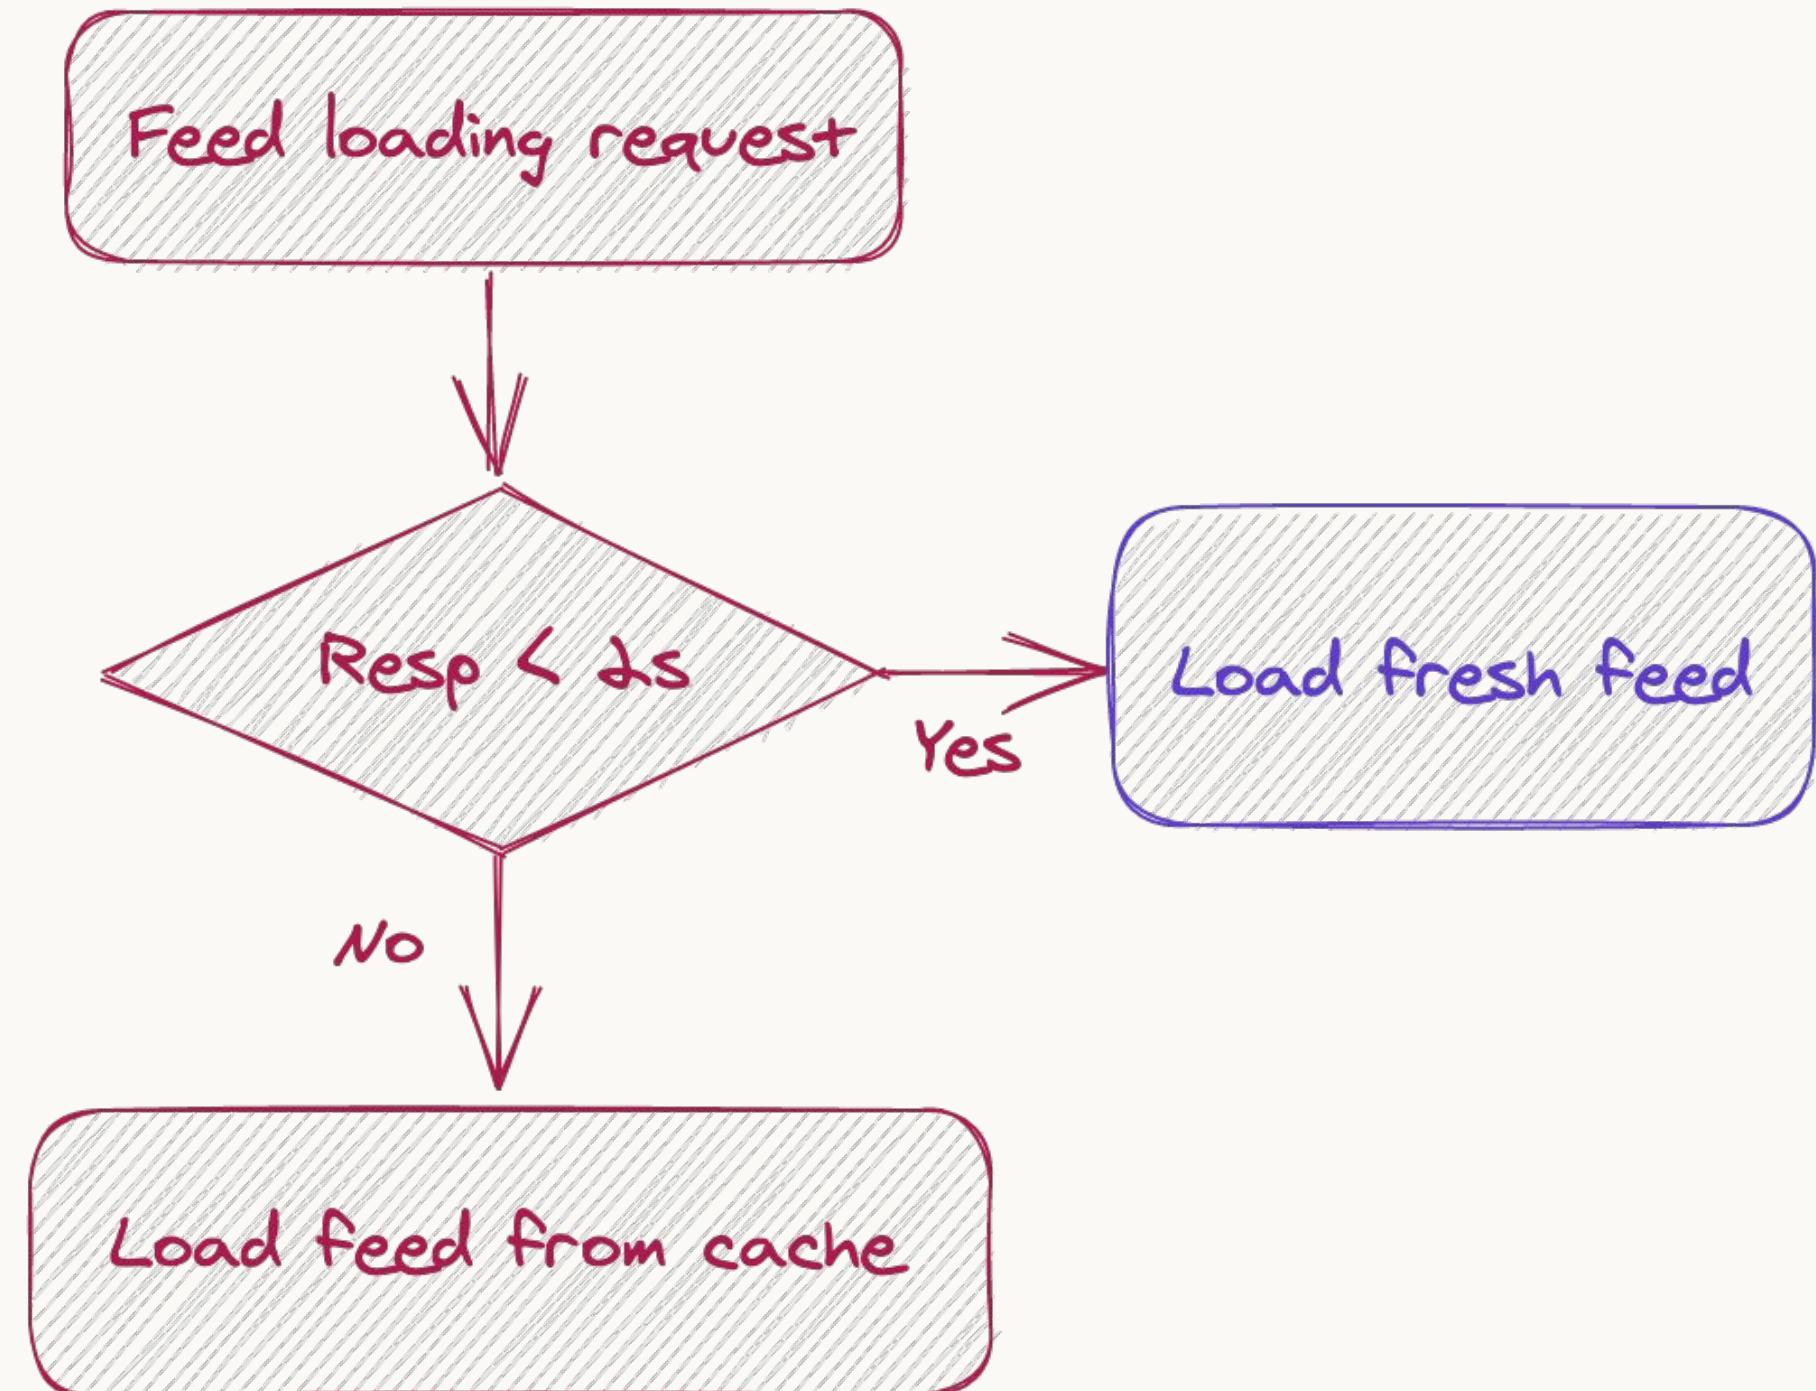
Flow control > congestion window

2MB



Clever heuristics aren't always that clever

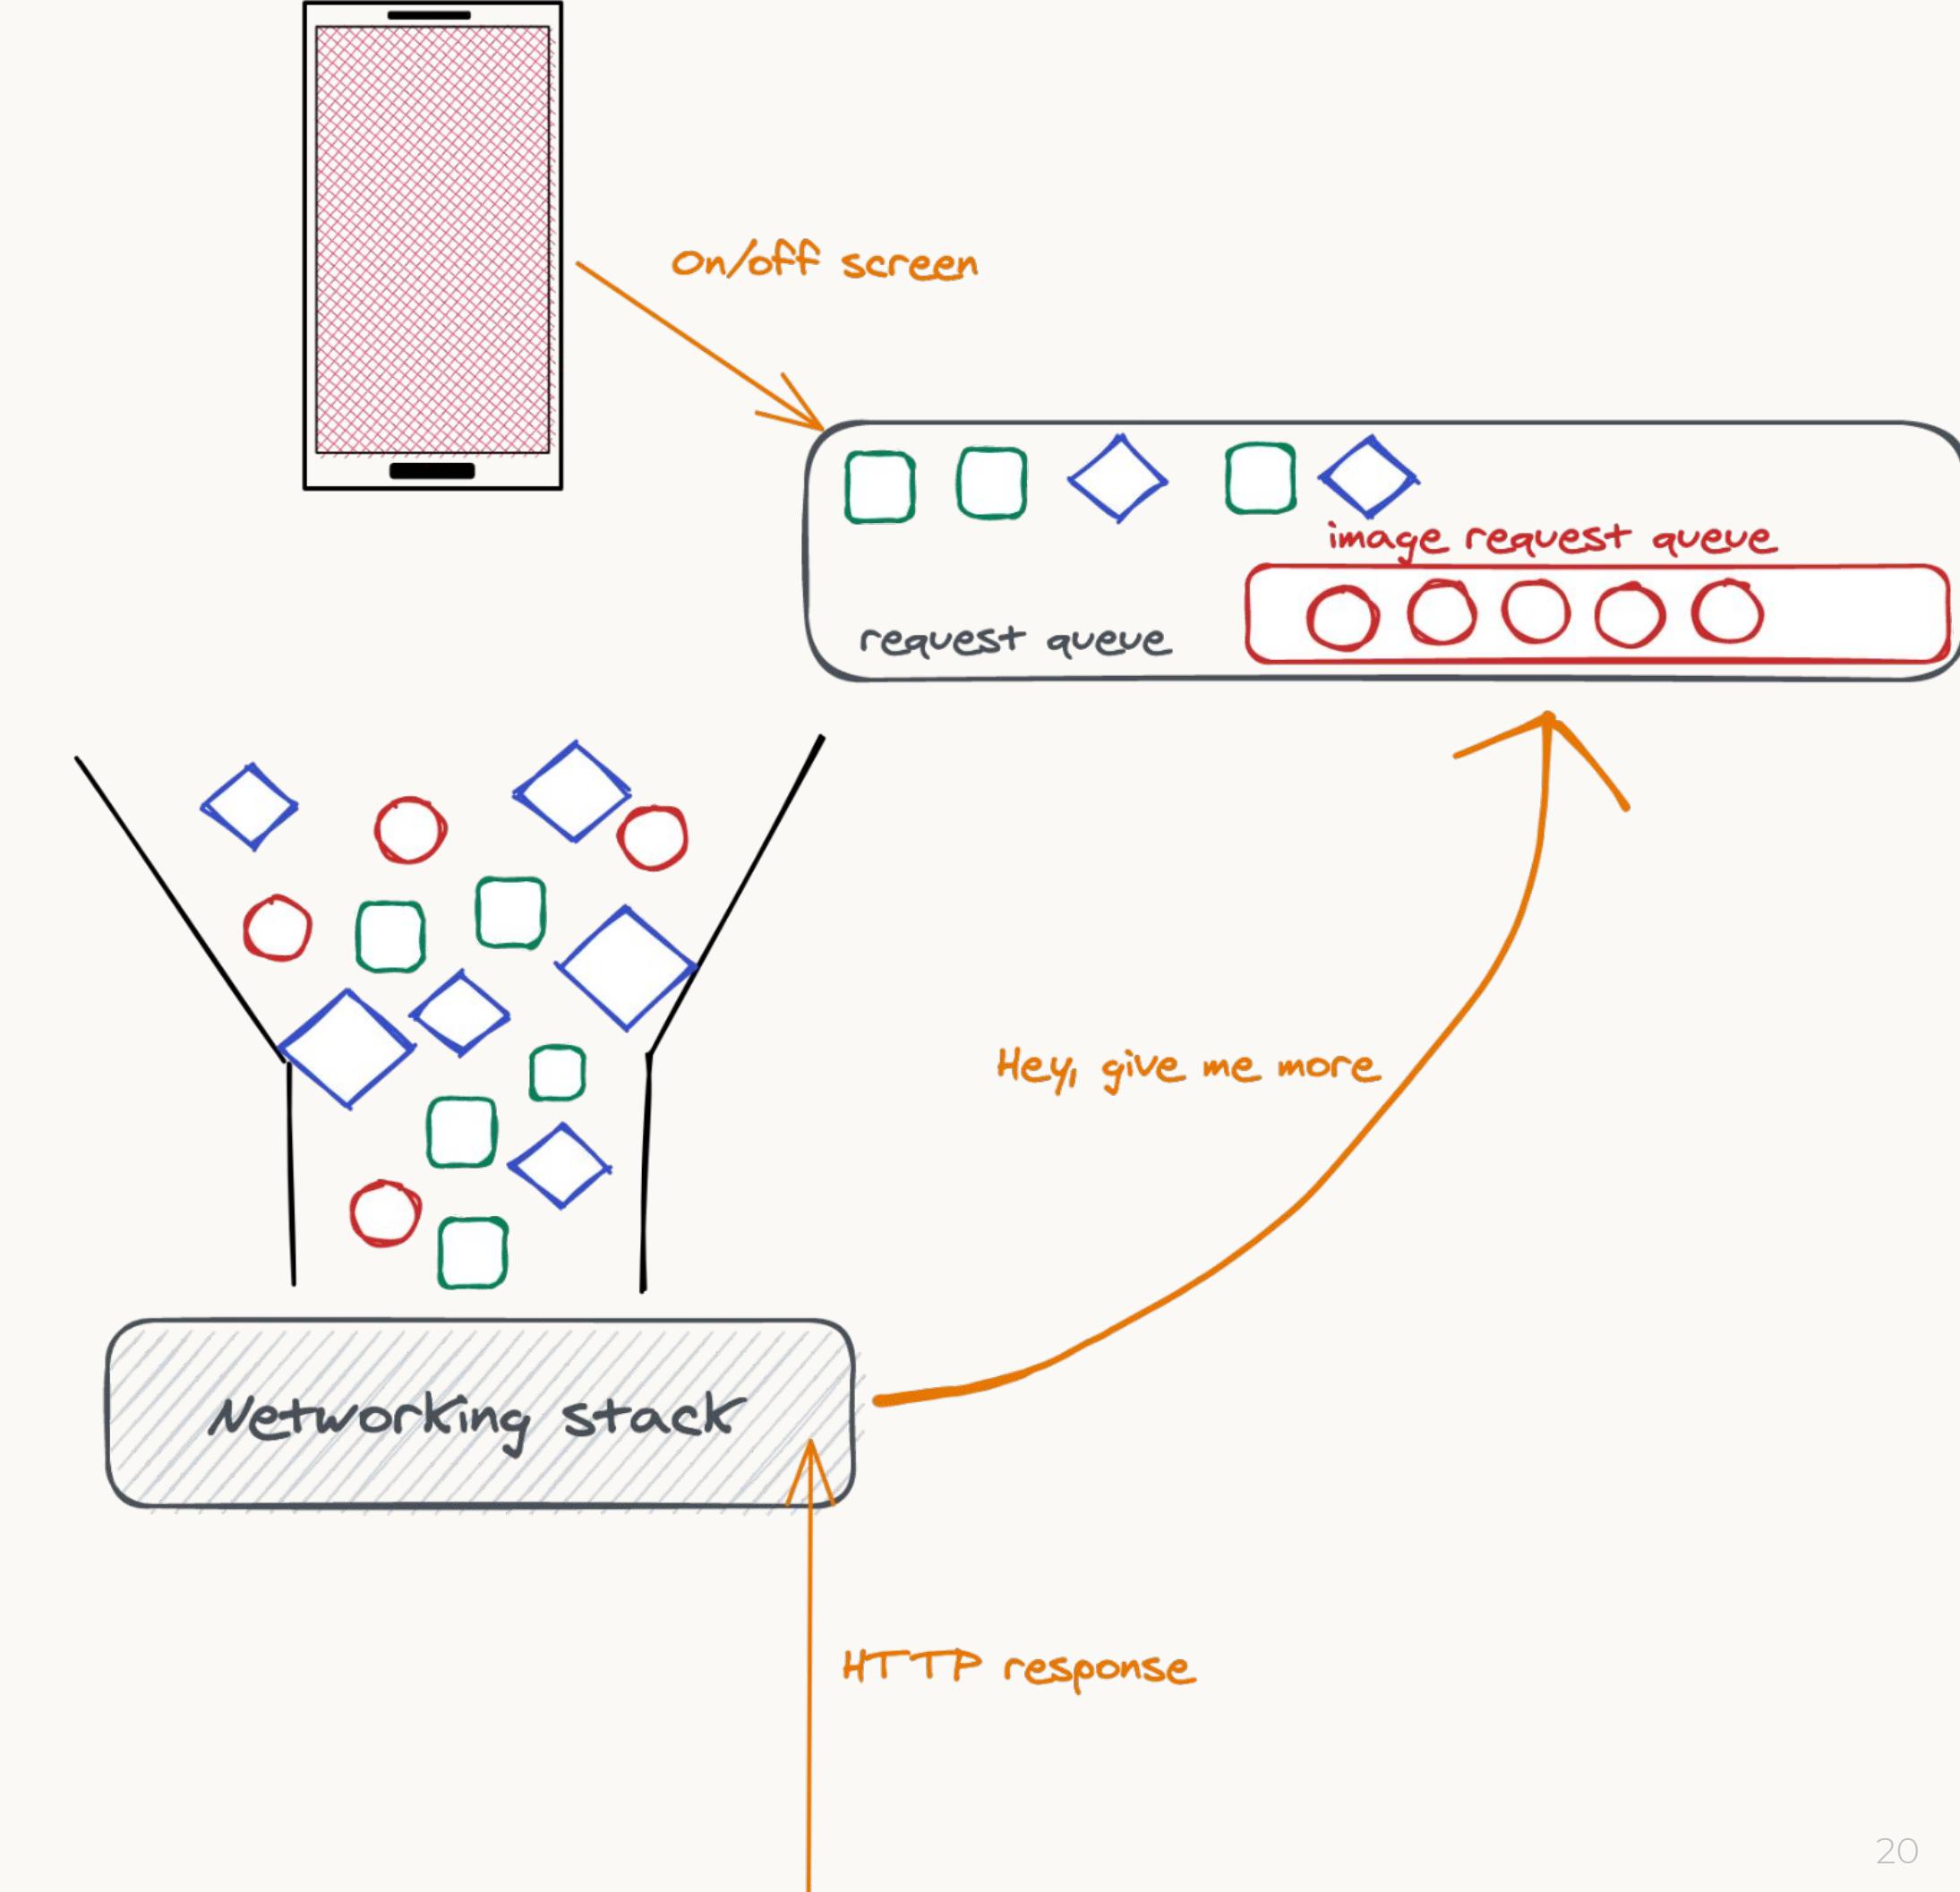
- Enabled QUIC and HTTP/3 for GraphQL
 - -> Regressed images and video



App request queue

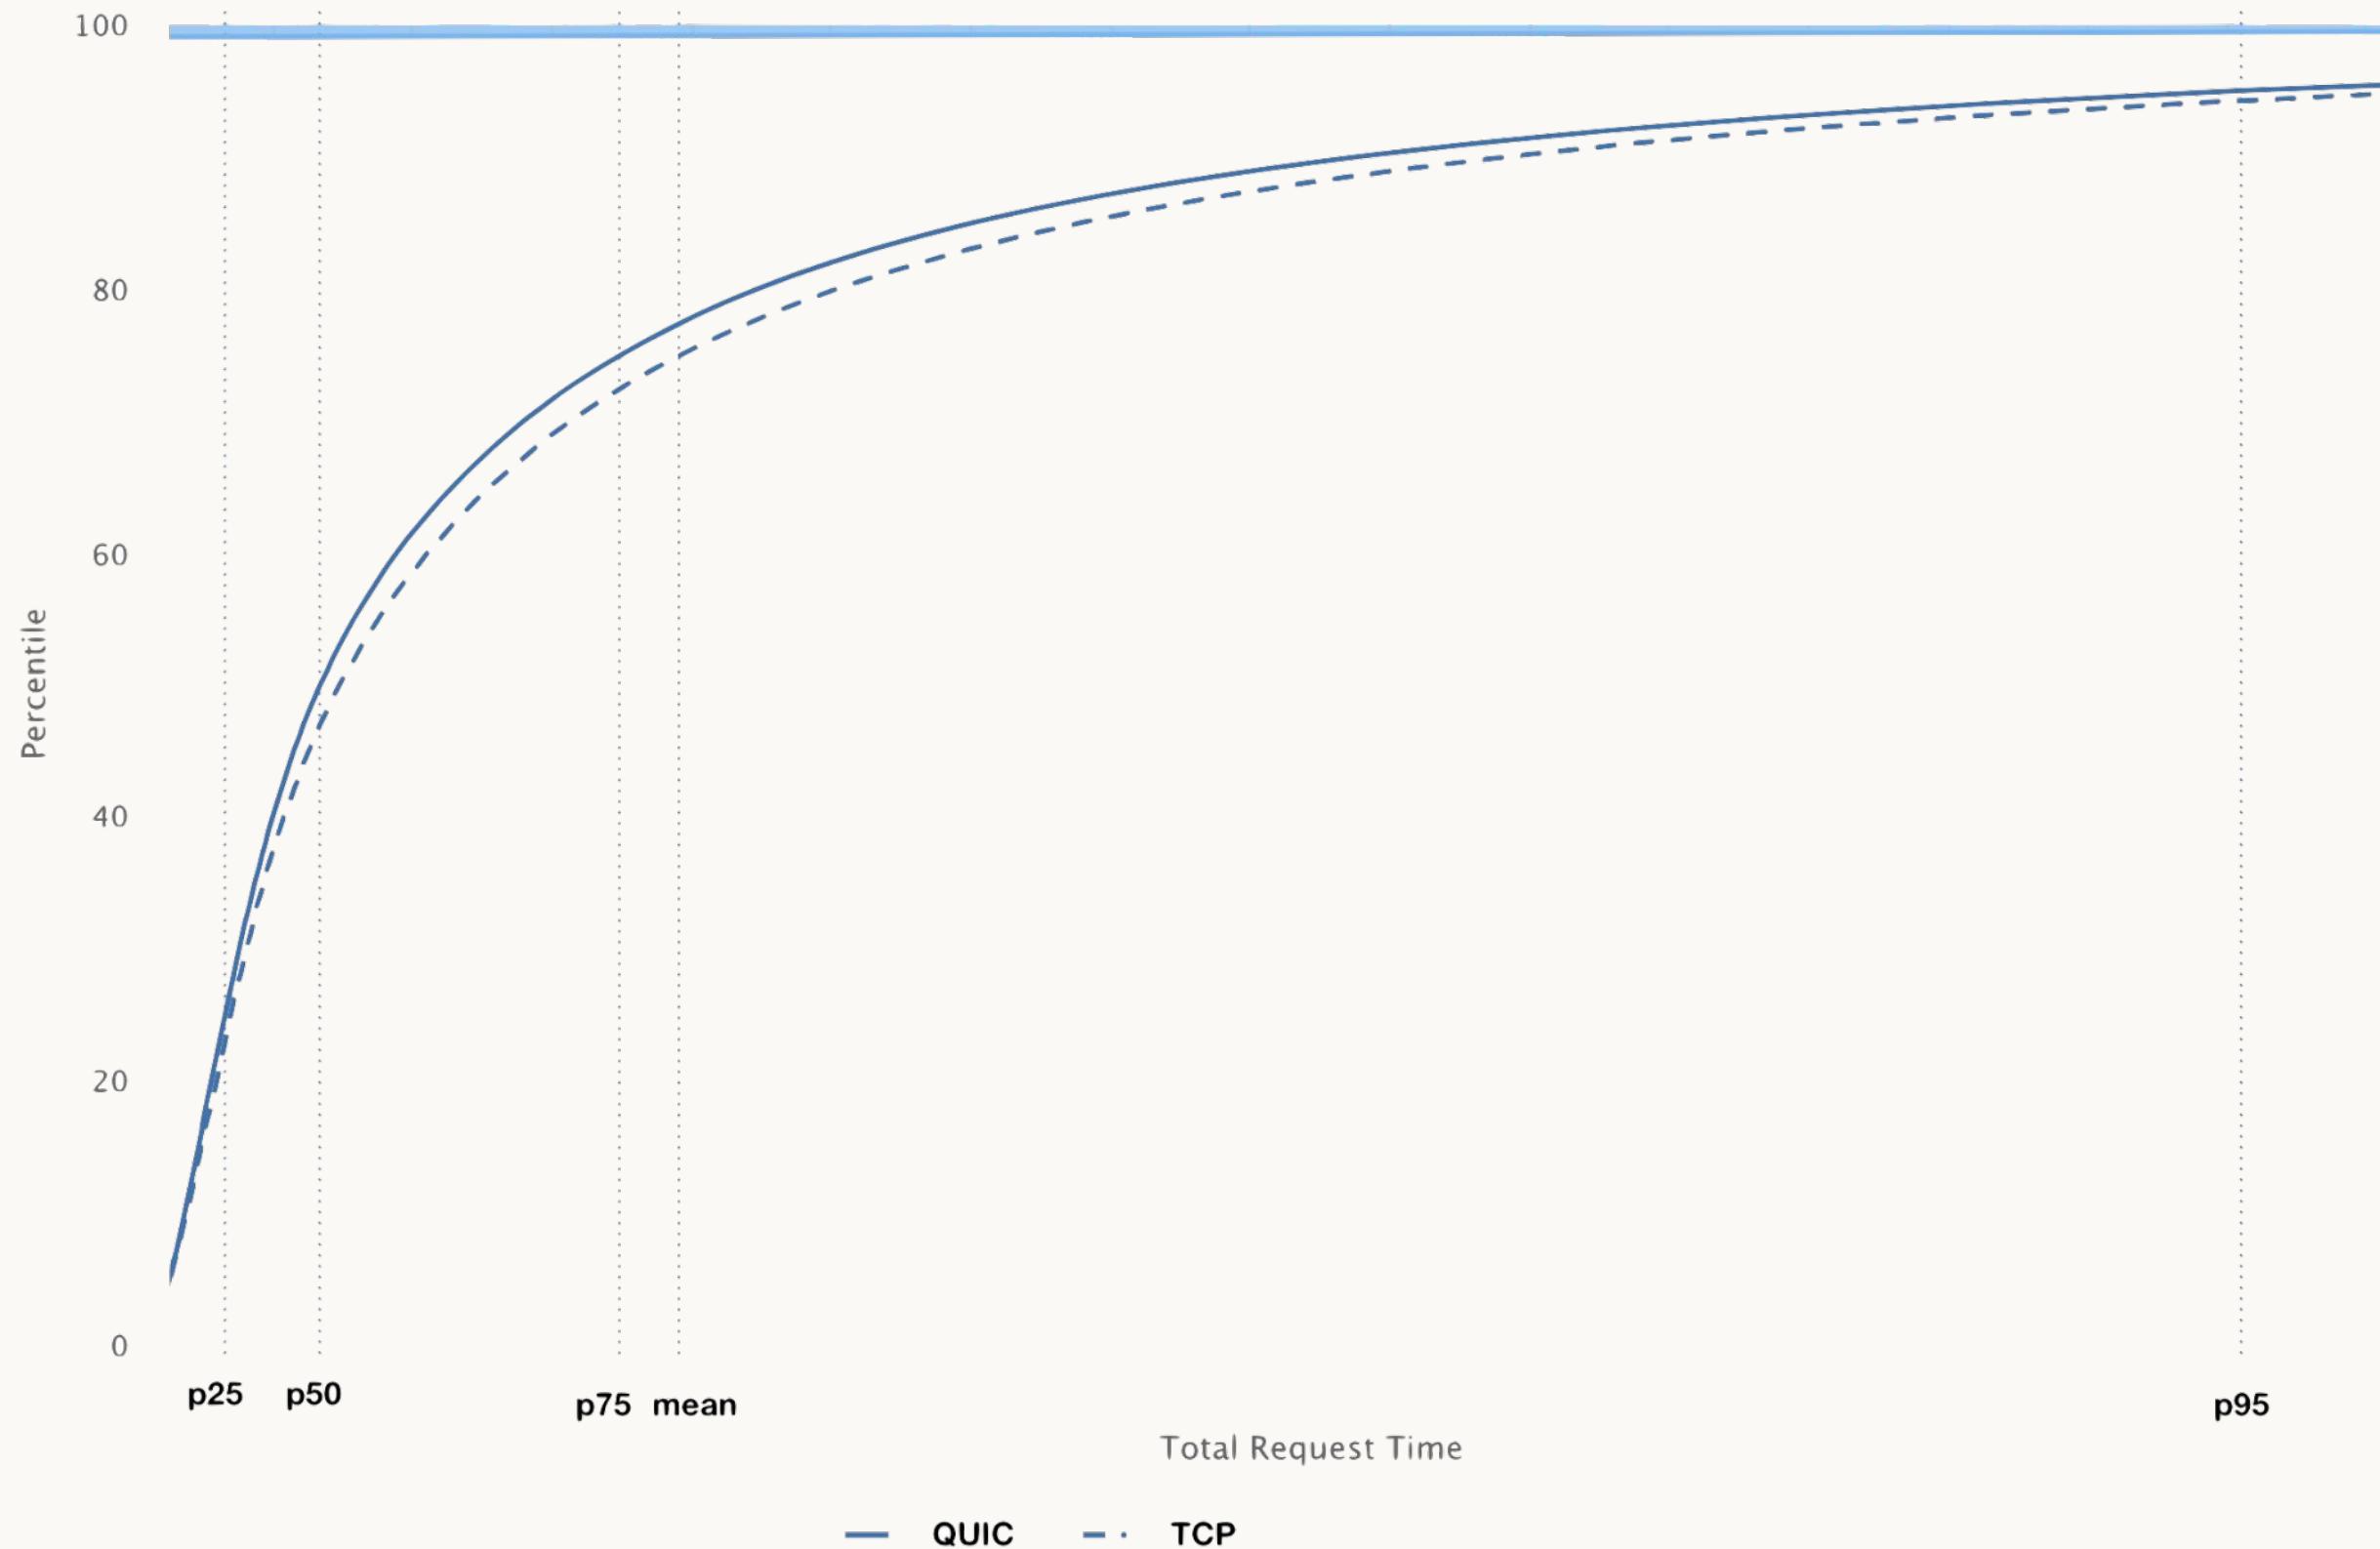
- Higher video quality over QUIC
 - -> Longer image queueing time.

= API request
 = image request
 = video request

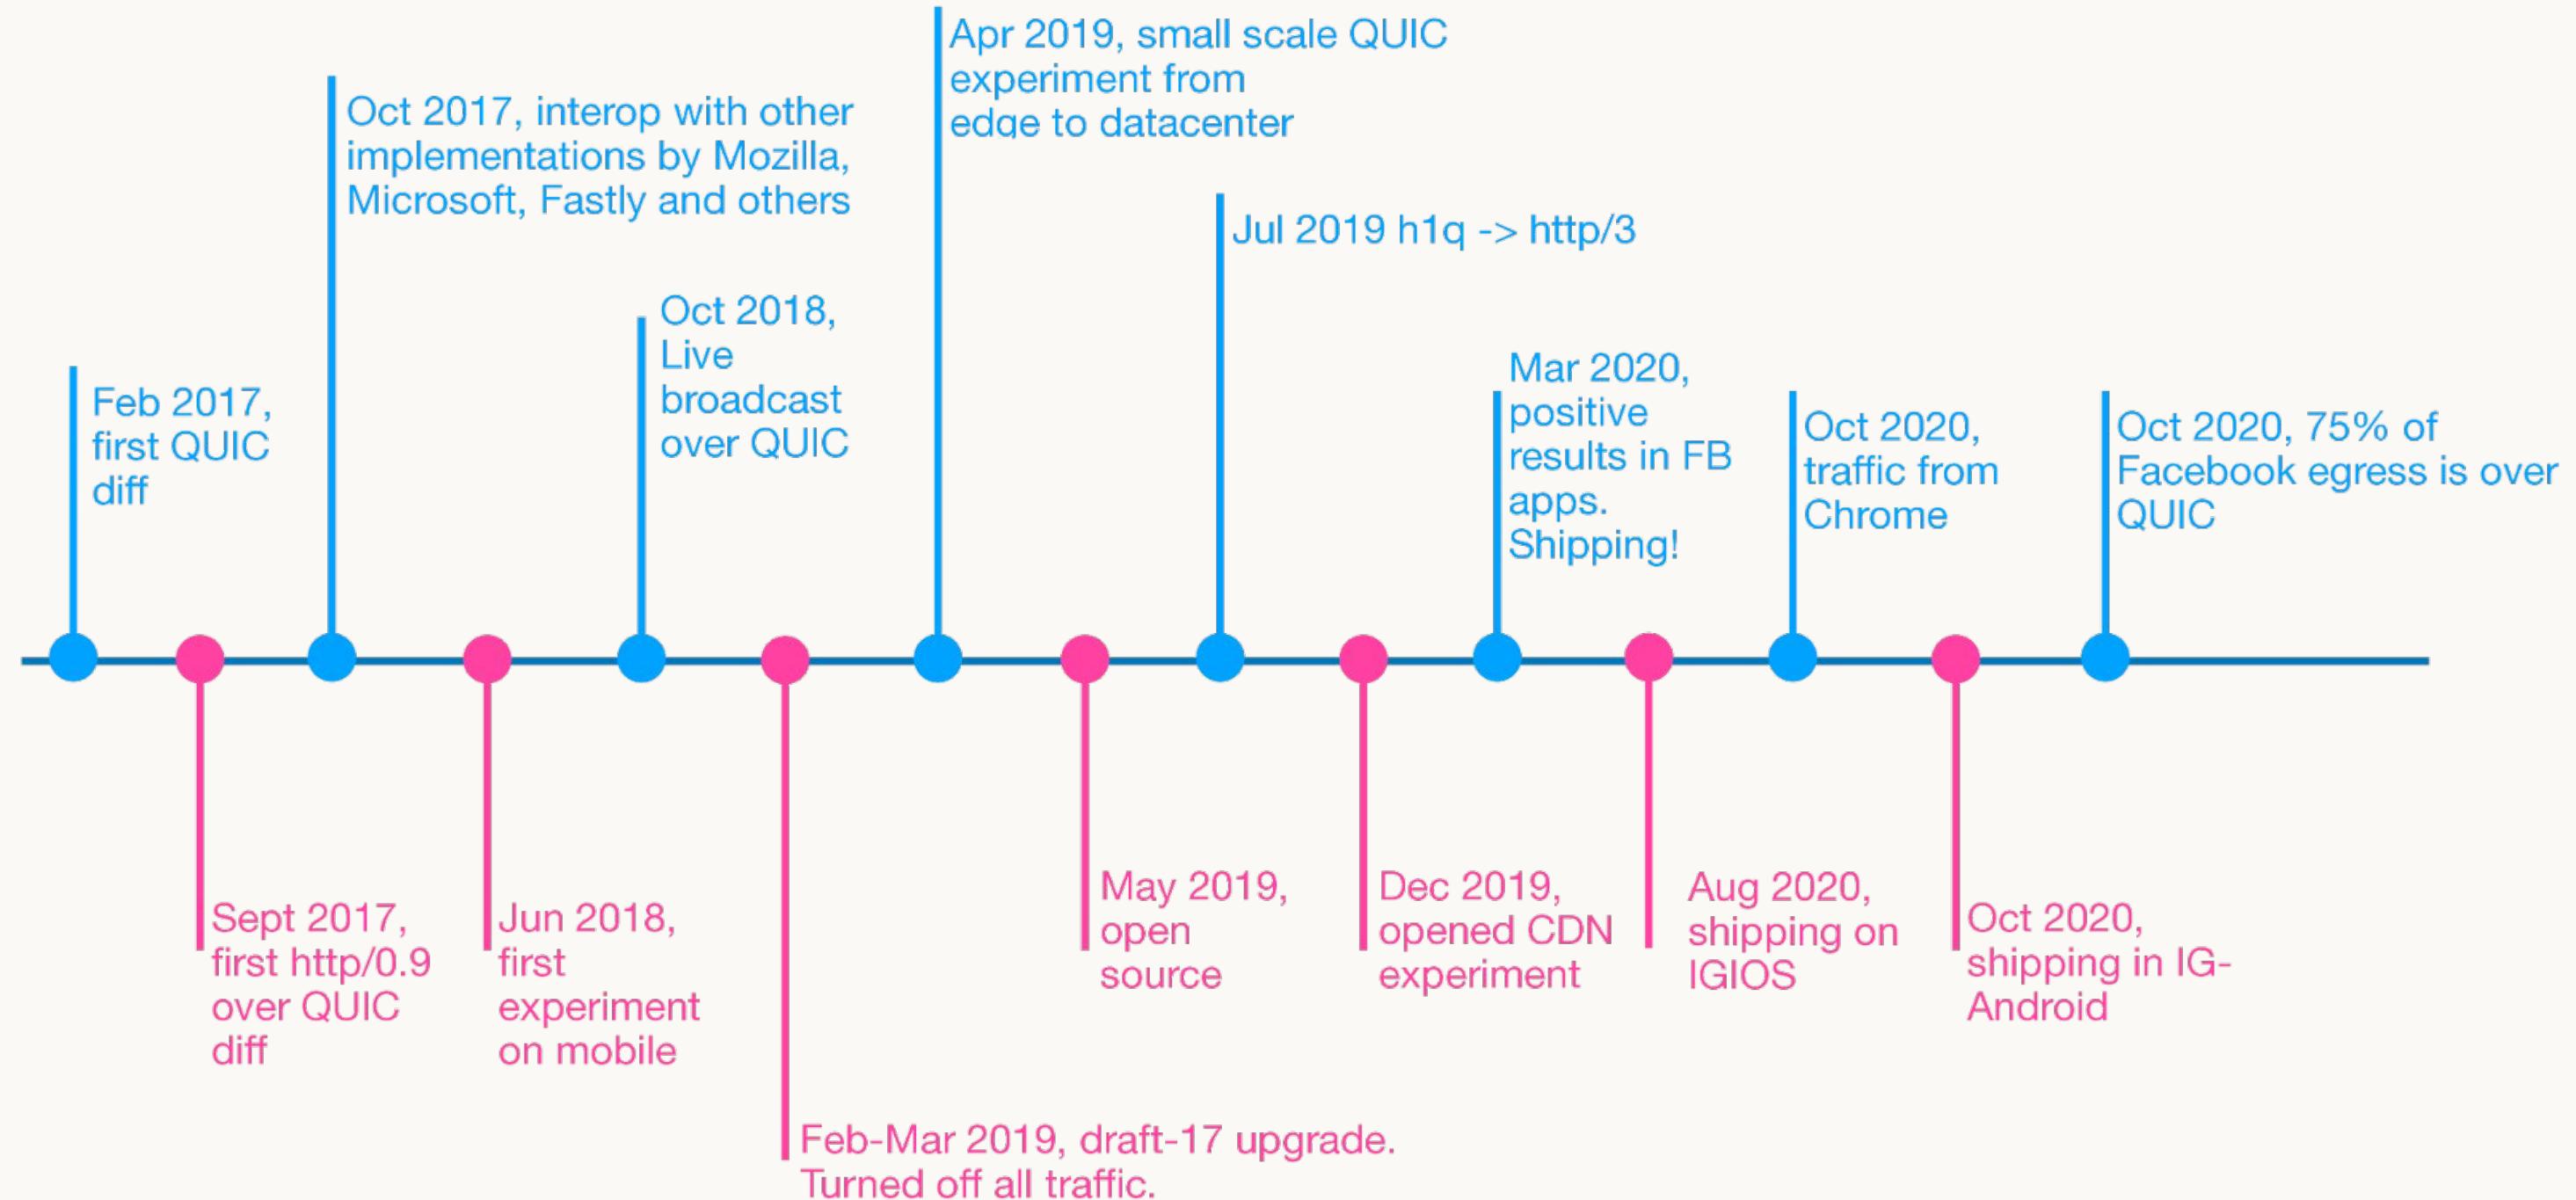


QUIC Had a Transformative Effect on Video

- 22% increase in MTBR
- 20% decrease in video stalls
- 8% reduced video error rate



The voyage



Takeaways

Tuning is needed for QUIC to deliver great result

- Application behavior can be tricky and implicitly coupled to TCP behavior.
- Often HTTP and Transport metrics look good while Application metrics will regress.
- Putting in the effort to tune usage for QUIC is worth the effort.

Future Work

0-RTT and Connection migration

- Didn't need them to show improvements over TCP.
- Connection reuse from pooling already extremely high in our stack.
- 0-RTT in progress.
- Connection migration has a lot of operational complexity.



Thank you